
DPC++ Reference

Intel

Apr 25, 2020

CONTENTS

| | | |
|----------|---------------------------------------------------------------|-----------|
| 1 | Introduction | 1 |
| 1.1 | Overview | 1 |
| 1.2 | Data Parallel C++ (DPC++) | 2 |
| 1.3 | Related Documentation | 3 |
| 2 | Programming Model | 5 |
| 2.1 | Sample Program | 5 |
| 2.2 | Platform Model | 9 |
| 2.3 | Execution Model | 11 |
| 2.4 | Memory Model | 13 |
| 2.4.1 | Memory Objects | 15 |
| 2.4.2 | Accessors | 15 |
| 2.4.3 | Synchronization | 15 |
| 2.4.4 | Unified Shared Memory | 16 |
| 2.4.5 | Memory Types | 16 |
| 2.5 | Kernel Programming Model | 17 |
| 2.5.1 | C++ Language Requirements | 17 |
| 2.5.2 | Error Handling | 19 |
| 2.5.3 | Fall Back | 19 |
| 3 | Language | 21 |
| 3.1 | Keywords | 22 |
| 3.2 | Preprocessor Directives and Macros | 22 |
| 3.3 | Standard Library Classes Required for the Interface | 22 |
| 4 | Programming Interface | 23 |
| 4.1 | Header File | 23 |
| 4.2 | Namespaces | 23 |
| 4.3 | Class Descriptions | 23 |
| 4.3.1 | Runtime Classes | 23 |
| 4.3.1.1 | device_selector | 23 |
| 4.3.1.2 | platform | 24 |
| 4.3.1.3 | context | 28 |
| 4.3.1.4 | device | 30 |
| 4.3.1.5 | queue | 33 |
| 4.3.1.6 | event | 35 |
| 4.3.2 | Data Access | 37 |
| 4.3.2.1 | accessor | 37 |
| 4.3.2.2 | Atomic | 39 |
| 4.3.2.3 | Buffer | 40 |

| | | |
|----------|--------------------------------------|-----------|
| 4.3.2.4 | Image | 43 |
| 4.3.2.5 | Address Spaces | 44 |
| 4.3.2.6 | Samplers | 46 |
| 4.3.3 | Expressing Parallelism | 46 |
| 4.3.3.1 | Range | 46 |
| 4.3.3.2 | Nd_range | 49 |
| 4.3.3.3 | ID | 49 |
| 4.3.3.4 | Item | 51 |
| 4.3.3.5 | Nd_item | 52 |
| 4.3.3.6 | Group | 54 |
| 4.3.3.7 | Device Event | 55 |
| 4.3.3.8 | Command Group Handler | 55 |
| 4.3.3.9 | Kernel | 58 |
| 4.3.3.10 | Program | 59 |
| 4.3.4 | Error Handling | 60 |
| 4.3.4.1 | Exception | 60 |
| 4.3.5 | Stream | 61 |
| 4.3.6 | Vec and Swizzled Vec | 62 |
| 4.3.7 | Built-in Types & Functions | 66 |
| 4.3.8 | Property Interface | 67 |
| 4.3.9 | Version | 67 |
| 5 | Glossary | 69 |
| 6 | Notices and Disclaimers | 71 |
| | Index | 73 |

INTRODUCTION

Obtaining high compute performance on today's modern computer architectures requires code that is optimized, power efficient, and scalable. The demand for high performance continues to increase due to needs in AI, video analytics, data analytics, as well as in traditional high performance computing (HPC).

Modern workload diversity has resulted in a need for architectural diversity; no single architecture is best for every workload. A mix of scalar, vector, matrix, and spatial (SVMS) architectures deployed in CPU, GPU, AI, and FPGA *accelerators* is required to extract the needed performance.

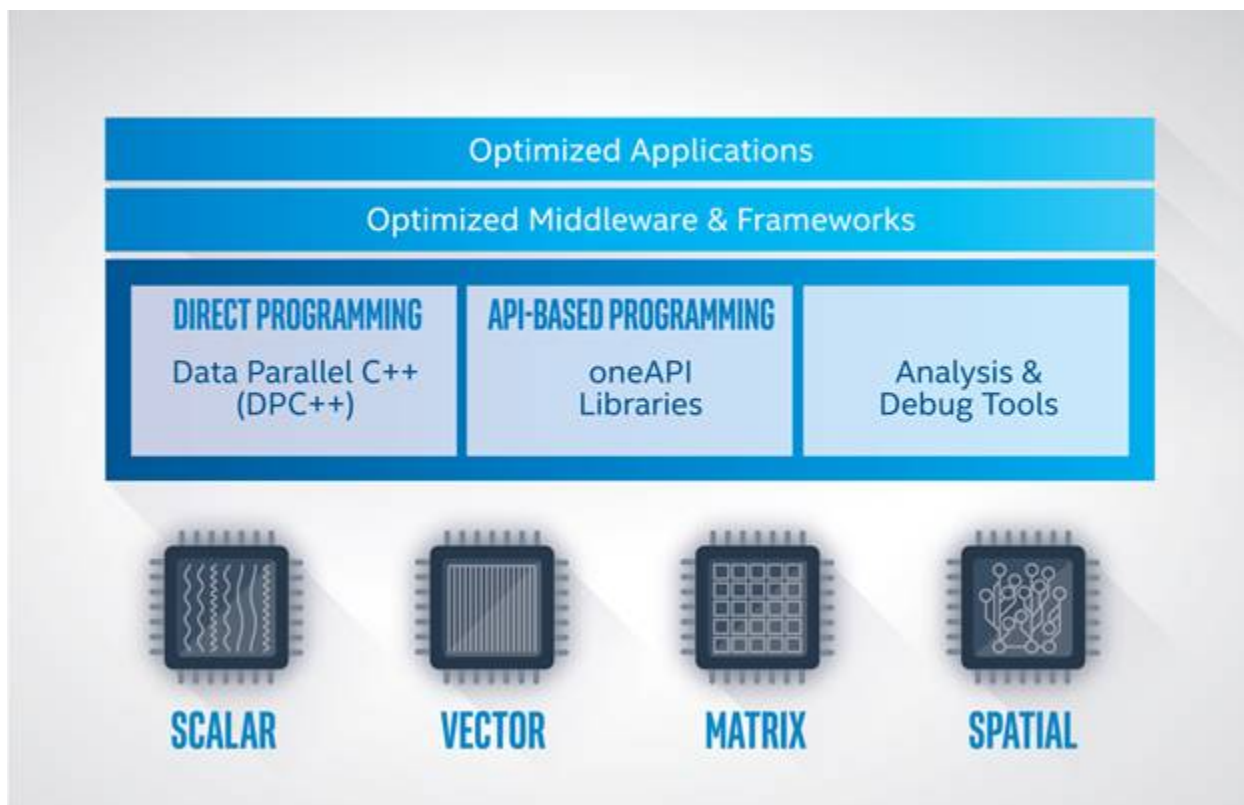
Today, coding for CPUs and accelerators requires different languages, libraries, and tools. That means each hardware platform requires completely separate software investments and provides limited application code reusability across different target architectures.

The oneAPI programming model simplifies the programming of CPUs and accelerators using modern C++ features to express parallelism with a programming language called Data Parallel C++ (DPC++). The DPC++ language enables code reuse for the host (such as a CPU) and accelerators (such as a GPU) using a single source language, with execution and memory dependencies clearly communicated. Mapping within the DPC++ code can be used to transition the application to run on the hardware, or set of hardware, that best accelerates the workload. A host is available to simplify development and debugging of device code, even on platforms that do not have an accelerator available.

Note: Not all programs can benefit from the single programming model offered by oneAPI. It is important to understand if your program can benefit and how to design, implement, and use the oneAPI programming model for your program.

1.1 Overview

The oneAPI programming model provides a comprehensive and unified portfolio of developer tools that can be used across hardware targets, including a range of performance libraries spanning several workload domains. The libraries include functions custom-coded for each target architecture, so the same function call delivers optimized performance across supported architectures. DPC++ is based on industry standards and open specifications to encourage ecosystem collaboration and innovation.



As shown in the figure above, applications that take advantage of the oneAPI programming model can execute on multiple target hardware platforms ranging from CPU to FPGA.

1.2 Data Parallel C++ (DPC++)

Data Parallel C++ (DPC++) is a high-level language designed for data parallel programming productivity. It is based on C++ for broad compatibility and uses common, familiar C and C++ constructs. The language seeks to deliver performance on par with other compiled languages, such as standard C++ compiled code, and uses C++ class libraries to allow the compiler to interpret the code and run on various supported architectures.

DPC++ is based on SYCL* from the Khronos* Group to support data parallelism and heterogeneous programming. In addition, Intel is pursuing extensions to SYCL with the aim of providing value to customer code and working with the standards organization for adoption. For instance, the DPC++ language includes an implementation of unified shared memory to ease memory usage between the host and the accelerators. These features are being driven into a future version of the SYCL language. For more details about SYCL, refer to version 1.2.1 of the [SYCL Specification](#).

While DPC++ applications can run on any supported target hardware, tuning is required to gain the best performance advantage on a given target architecture. For example, code tuned for a CPU likely will not run as fast on a GPU accelerator without modification. This guide aims to help developers understand how to program using the oneAPI programming model and how to target and optimize for the appropriate architecture to achieve optimal application performance.

1.3 Related Documentation

The following documents are useful starting points for developers getting started with oneAPI projects. This document assumes you already have a basic understanding of the oneAPI programming model concepts.

[Get Started with Intel oneAPI for Linux*](#)

[Get Started with Intel oneAPI for Windows*](#)

[Intel oneAPI Base Toolkit Release Notes](#)

[SYCL* Specification \(for version 1.2.1\)](#)

PROGRAMMING MODEL

The oneAPI programming model is based upon the [SYCL* Specification](#). The specification presents a general heterogeneous compute capability by detailing four models. These models categorize the actions a developer needs to perform to employ one or more devices as an accelerator. Aspects of these models appear in every program that employs the oneAPI programming model. These models are summarized as:

- *Platform Model*: Specifies the *host* and *device*.
- *Execution Model*: Specifies the *command queue* and issuing commands for execution on the device(s).
- *Memory Model*: Defines how the host and devices interact with memory.
- *Kernel Programming Model*: Defines the code that executes on the device(s). This code is known as the kernel.

The programming language for oneAPI is Data Parallel C++ (DPC++) and employs modern features of the C++ language to enact its parallelism. In fact, when writing programs that employ the oneAPI programming model, the programmer routinely uses language features such as C++ lambdas, templates, `parallel_for`, and closures.

Tip

If you are unfamiliar with these C++11 and later language features, consult other C++ language references and gain a basic understanding before continuing.

When evaluating and learning oneAPI, keep in mind that the programming model is general enough to accommodate multiple classes of accelerators; therefore, there may be a greater number of API calls required to access the accelerators than more constrained APIs, such as those only accessing one type of accelerator.

One of the primary motivations for DPC++ is to provide a higher-level programming language than OpenCL™ C code, which it is based upon. Readers familiar with OpenCL programs will see many similarities to and differences from OpenCL code. This chapter points out similarities and differences where appropriate. This chapter also points to portions of the SYCL Specification for further information.

2.1 Sample Program

The following code sample contains a program that employs the oneAPI programming model to compute a vector addition. The program computes the formula $c = a + b$ across arrays, *a* and *b*, each containing 1024 elements, and places the result in array *c*. The following discussion focuses on sections of code identified by line number in the sample. The intent with this discussion is to highlight the required functionality inherent when employing the programming model.

Note: Keep in mind that this sample code is intended to illustrate the four models that comprise the oneAPI programming model; it is not intended to be a typical program or the simplest in nature.

```

1  #include <vector>
2  #include <CL/sycl.hpp>
3
4  #define SIZE 1024
5
6  namespace sycl = cl::sycl;
7
8  void show_platforms() {
9      auto platforms = sycl::platform::get_platforms();
10
11     for (auto &platform : platforms) {
12         std::cout << "Platform: "
13                 << platform.get_info<sycl::info::platform::name>()
14                 << std::endl;
15
16         auto devices = platform.get_devices();
17         for (auto &device : devices) {
18             std::cout << "  Device: "
19                     << device.get_info<sycl::info::device::name>()
20                     << std::endl;
21         }
22     }
23 }
24
25 void vec_add(int *a, int *b, int *c) {
26     sycl::range<1> a_size{SIZE};
27
28     sycl::default_selector device_selector;
29     sycl::queue q(device_selector);
30
31     sycl::buffer<int, 1> a_device(a, a_size);
32     sycl::buffer<int, 1> b_device(b, a_size);
33     sycl::buffer<int, 1> c_device(c, a_size);
34
35     q.submit([&](sycl::handler &cgh) {
36         auto c_res = c_device.get_access<sycl::access::mode::write>(cgh);
37         auto a_in = a_device.get_access<sycl::access::mode::read>(cgh);
38         auto b_in = b_device.get_access<sycl::access::mode::read>(cgh);
39
40         cgh.parallel_for<class ex1>(a_size, [=](sycl::id<1> idx) {
41             c_res[idx] = a_in[idx] + b_in[idx];
42         });
43     });
44 }
45
46 int main() {
47     int a[SIZE], b[SIZE], c[SIZE];
48
49     for (int i = 0; i<SIZE; ++i) {
50         a[i] = i;
51         b[i] = -i;
52         c[i] = i;
53     }
54
55     show_platforms();
56     vec_add(a, b, c);
57

```

(continues on next page)

(continued from previous page)

```

58     return 0;
59 }

```

A DPC++ program has the *single source* property, which means the *host code* and the *device code* can be placed in the same file so that the compiler treats them as the same compilation unit. This can potentially result in performance optimizations across the boundary between host and device code. The single source property differs from a programming model like OpenCL software technology where the host code and device code are typically in different files, and the host and device compiler are different entities, which means no optimization can occur between the host and device code boundary. Therefore, when scrutinizing a DPC++ program, the first step is to understand the delineation between host code and device code. To be more specific, DPC++ programs are delineated into different scopes similar to programming language scope, which is typically expressed via { and } in many languages.

The three types of scope in a DPC++ program include:

- *application scope*: Code that executes on the host
- *command group scope*: Code that acts as the interface between the host and device
- *kernel scope*: Code that executes on the device

In this example, command group scope comprises lines 45 through 54 and includes coordination and data passing operations required in the program to enact control and communication between the host and the device.

```

45     d_queue.submit([&](sycl::handler &cgh) {
46         auto c_res = c_device.get_access<sycl::access::mode::write>(cgh);
47         auto a_in = a_device.get_access<sycl::access::mode::read>(cgh);
48         auto b_in = b_device.get_access<sycl::access::mode::read>(cgh);
49
50         cgh.parallel_for<class ex1>(a_size, [=](sycl::id<1> idx) {
51             c_res[idx] = a_in[idx] + b_in[idx];
52         });
53
54     });

```

Kernel scope, which is nested in the command group scope, comprises lines 50 to 52. Application scope consists of all the other lines not in command group or kernel scope. Syntactically, definitions are included from the top level include file; `sycl.hpp` and namespace declarations can be added for convenience.

The function of each of the lines and its classification into one of the four models are detailed as follows:

- Lines 2 and 6 – `include` and `namespace` – programs employing the oneAPI programming model require the include of `cl/sycl.hpp`. It is recommended to employ the `namespace` statement at line 6 to save typing repeated references into the `cl::sycl` namespace.

```

2     #include <CL/sycl.hpp>
3
4     #define SIZE 1024
5
6     namespace sycl = cl::sycl;

```

- Lines 20 to 36 – Platform model – programs employing the oneAPI programming model can query the host for available platforms and can either select one to employ for execution or allow the oneAPI runtime to choose a default platform. A platform defines the relationship between the host and device(s). The platform may have a number of devices associated with it and a program can specify which device(s) to employ for execution or allow the oneAPI runtime to choose a default device.

```

20     auto platforms = sycl::platform::get_platforms();
21
22     for (auto &platform : platforms) {
23
24         std::cout << "Platform: "
25             << platform.get_info<sycl::info::platform::name>()
26             << std::endl;
27
28
29         auto devices = platform.get_devices();
30         for (auto &device : devices) {
31             std::cout << "  Device: "
32                 << device.get_info<sycl::info::device::name>()
33                 << std::endl;
34         }
35     }
36 }

```

- Lines 39 and 45 – Execution model – programs employing the oneAPI programming model define command queues that issue command groups. The command groups control execution on the device.

```

39     sycl::queue d_queue(device_selector);
40
41     sycl::buffer<int, 1>  a_device(a.data(), a_size);
42     sycl::buffer<int, 1>  b_device(b.data(), a_size);
43     sycl::buffer<int, 1>  c_device(c.data(), a_size);
44
45     d_queue.submit([&](sycl::handler &cgh) {

```

- Lines 41 to 43 and lines 46 to 48 – Memory model – programs employing the oneAPI programming model may use buffers and accessors to manage memory access between the host and devices. In this example, the arrays, a, b, and c are defined and allocated on the host. Buffers, a_device, b_device, and c_device, are declared to hold the values from a, b, and c respectively so the device can compute the vector addition. The accessors, a_in and b_in, denote that a_device and b_device are to have read only access on the device. The accessor c_res denotes that c_device is to allow write access from the device.

```

41     sycl::buffer<int, 1>  a_device(a.data(), a_size);
42     sycl::buffer<int, 1>  b_device(b.data(), a_size);
43     sycl::buffer<int, 1>  c_device(c.data(), a_size);
44
45     d_queue.submit([&](sycl::handler &cgh) {
46         auto c_res = c_device.get_access<sycl::access::mode::write>(cgh);
47         auto a_in = a_device.get_access<sycl::access::mode::read>(cgh);
48         auto b_in = b_device.get_access<sycl::access::mode::read>(cgh);

```

- Line 50 to 52 – Kernel Programming Model – The C++ language `parallel_for` statement denotes that the code enclosed in its scope will execute in parallel across the *processing elements* of the device. This example code employs a C++ lambda to represent the kernel.

```

50         cgh.parallel_for<class ex1>(a_size, [=](sycl::id<1> idx) {
51             c_res[idx] = a_in[idx] + b_in[idx];
52         });

```

- Line 17 and 56 – Scope and Synchronization – Memory operations between the buffers and actual host memory

execute in an asynchronous fashion. To ensure synchronization, the command queue is placed inside another scope at line 17 and 56 which tells the runtime to synchronize before the scope is exited as part of the buffer's destructors being executed. This practice is used in many programs.

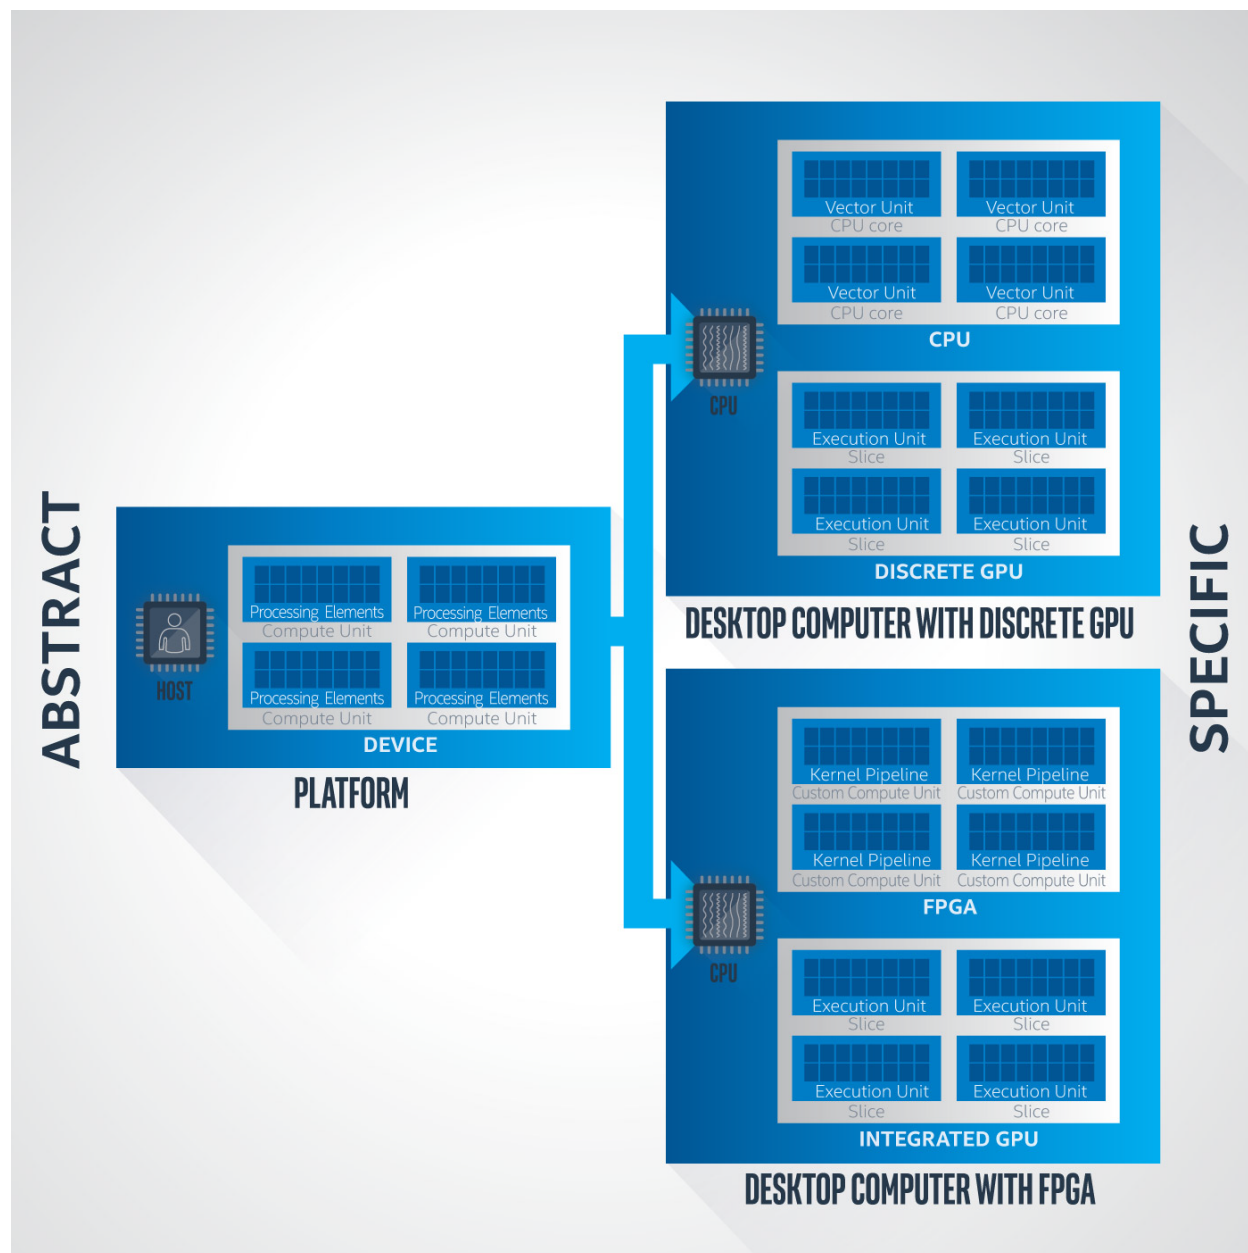
When compiled and executed, the sample program computes the 1024 element vector add in parallel on the accelerator. This assumes the accelerator has multiple compute elements capable of executing in parallel. This sample illustrates the models that the software developer will need to employ in their program. The next sections discuss in more details those four models: the Platform model, the Execution model, the Memory model, and the Kernel model.

2.2 Platform Model

The platform model for oneAPI is based upon the SYCL* platform model. It specifies a host controlling one or more devices. A host is the computer, typically a CPU-based system executing the primary portion of a program, specifically the application scope and the command group scope. The host coordinates and controls the compute work that is performed on the devices. A device is an accelerator, a specialized component containing compute resources that can quickly execute a subset of operations typically more efficiently than the CPUs in the system. Each device contains one or more compute units that can execute several operations in parallel. Each compute unit contains one or more *processing elements* that serve as the individual engine for computation.

A system can instantiate and execute several platforms simultaneously, which is desirable because a particular platform may only target a subset of the available hardware resources on a system. However, in the typical case, a system will have one platform comprised of one or more supported devices, and the compute resources made available by those devices.

The following figure provides a visual depiction of the relationships in the platform model. One host communicates with one or more devices. Each device can contain one or more compute units. Each compute unit can contain one or more processing elements.



The platform model is general enough to be mapped to several different types of devices and lends to the functional portability of the programming model. The hierarchy on the device is also general and can be mapped to several different types of accelerators from FPGAs to GPUs and ASICs as long as these devices support the minimal requirements of the oneAPI programming model. Consult the [Intel oneAPI Base Toolkit System Requirements](#) for more information.

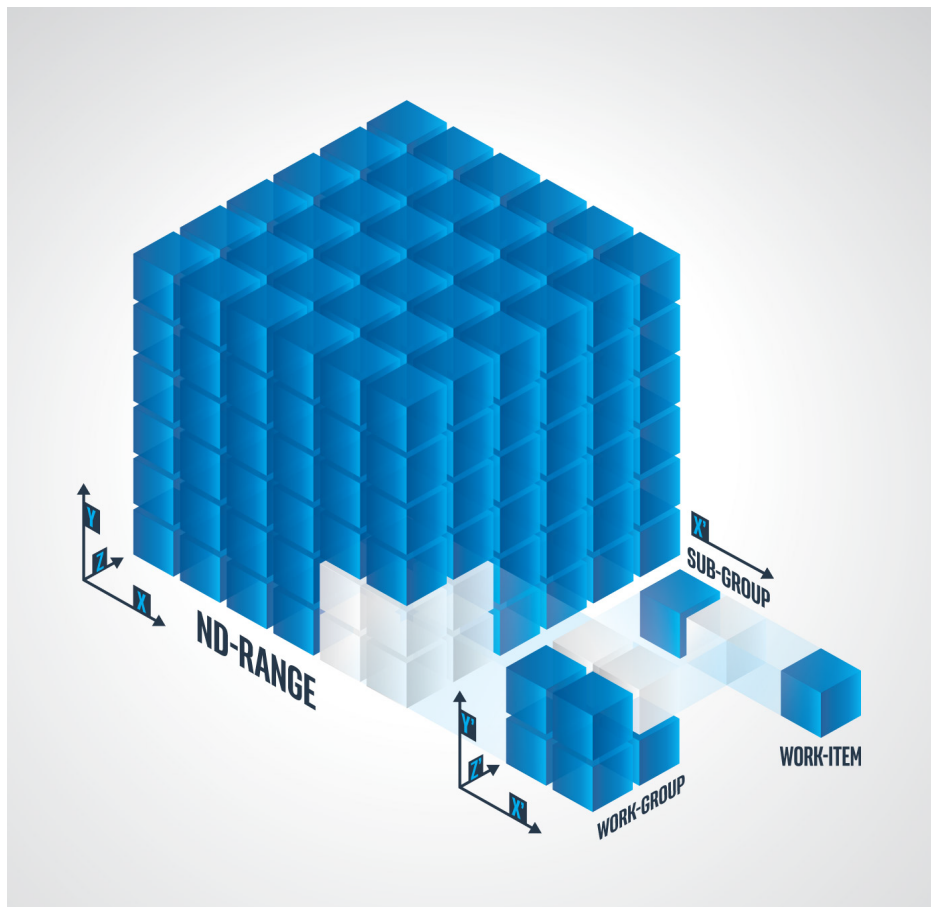
2.3 Execution Model

The execution model is based upon the SYCL* execution model. It defines and specifies how code, termed kernels, execute on the host and the devices.

The host execution model coordinates execution and data management between the host and devices via command groups. The command groups, which are groupings of commands like kernel invocation and accessors, are submitted to queues for execution. Accessors, which are formally part of the memory model, also communicate ordering requirements of execution. A program employing the execution model declares and instantiates queues. Queues can execute with an in-order or out-of-order policy controllable by the program. In-order execution is an Intel extension.

The device execution model specifies how computation is accomplished on the accelerator. Compute ranging from small one-dimensional data to large multidimensional data sets are allocated across a hierarchy of *nd-range*, *work-group*, *sub-group*, and *work-item*, which are all specified when the work is submitted to the command queue. It is important to note that the actual kernel code represents the work that is executed for one work-item. The code outside of the kernel controls just how much parallelism is executed; the amount and distribution of the work is controlled by specification of the sizes of the ND-range and work-group.

The following figure depicts the relationship between an ND-range, work-group, sub-group, and work-item. The total amount of work is specified by the ND-range size. The grouping of the work is specified by the work-group size. The example shows the ND-range size of $X * Y * Z$, work-group size of $X' * Y' * Z'$, and subgroup size of X' . Therefore, there are $X * Y * Z$ work-items. There are $(X * Y * Z) / (X' * Y' * Z')$ work-groups and $(X * Y * Z) / X'$ subgroups.



When kernels are executed, the location of a particular work-item in the larger ND-range, work-group, or sub-group is important. For example, if the work-item is assigned to compute on specific pieces of data, a method of specification is necessary. Unique identification of the work-item is provided via intrinsic functions such as those in the `nd_item` class (`global_id`, `work_group_id`, and `local_id`).

The following code sample launches a kernel and displays the relationships of the previously discussed ND-range, work-group, and work-item.

```
#include<CL/sycl.hpp>
#include<iostream>
#define N 6
#define M 2
using namespace cl::sycl;
int main()
{
    queue defaultqueue;
    buffer<int,2> buf(range<2>(N,N));
    defaultqueue.submit([&](handler &cgh) {
        auto bufacc = buf.get_access<access::mode::read_write>(cgh);
        cgh.parallel_for<class ndim>(nd_range<2>(range<2>(N,N),
            range<2>(M,M)), [=](nd_item<2> i) {
            id<2> ind = i.get_global_id();
            bufacc[ind[0]][ind[1]] = ind[0]+ind[1];
        });
    });
    auto bufacc1 = buf.get_access<access::mode::read>();
    for(int i = 0; i < N; i++){
        for(int j = 0; j < M; j++){
            std::cout<<bufacc1[i][j]<<"\t";
            std::cout<<"\n";
        }
    }
    return 0;
}
```

ND-Range Parallelism Example

The following discusses the relationships in the use of the ND-range in the previous code sample.

- At line 12 is the nd-range declaration. `nd_range<2>` specifies a two-dimensional index space.
- The global range is specified by the first argument, `range<2>(N,N)`, which specifies the overall index space as 2 dimensions with size N by N.
- The second argument, `range<2>(M,M)` specifies the local work-group range as 2 dimensions with size M by M.
- Line 13 employs `nd_item<2>` to reference each work-item in the ND-range, and calls `get_global_id` to determine the index in the global buffer, `bufacc`.

The sub-group is an extension to the SYCL execution model and sits hierarchically between the `work_group` and `work_item`. The `sub_group` was created to align with typical hardware resources that contain a vector unit to execute several similar operations in parallel and in lock step.

2.4 Memory Model

The memory model for oneAPI is based upon the SYCL* memory model. It defines how the host and devices interact with memory. It coordinates the allocation and management of memory between the host and devices. The memory model is an abstraction that aims to generalize across and be adaptable to the different possible host and device configurations. In this model, memory resides upon and is owned by either the host or the device and is specified by declaring a memory object. There are two different types of memory objects, *buffers* and *images*. Interaction of these memory objects between the host and device is accomplished via an *accessor*, which communicates the desired location of access, such as host or device, and the particular mode of access, such as read or write.

Consider a case where memory is allocated on the host through a traditional malloc call. Once the memory is allocated on the host, a buffer object is created, which enables the host allocated memory to be communicated to the device. The `buffer` class communicates the type and number of items of that type to be communicated to the device for computation. Once a buffer is created on the host, the type of access allowed on the device is communicated via an `accessor` object, which specifies the type of access to the buffer. The general steps are summarized as:

1. Instantiate a `buffer` or `image` object.

The host or device memory for the `buffer` or `image` is allocated as part of the instantiation or is wrapped around previously allocated memory on the host.

2. Instantiate an `accessor` object.

The `accessor` specifies the required location of access, such as host or device, and the particular mode of access, such as read or write. It represents dependencies between uses of memory objects.

The following code sample is exercising different memory objects and accessors.

```
#include <vector>
#include <CL/sycl.hpp>
namespace sycl = cl::sycl;

#define SIZE 64

int main() {
    std::array<int, SIZE> a, c;
    std::array<sycl::float4, SIZE> b;
    for (int i = 0; i < SIZE; ++i) {
        a[i] = i;
        b[i] = (float)-i;
        c[i] = i;
    }

    {
        sycl::range<1> a_size{SIZE};

        sycl::queue d_queue;

        sycl::buffer<int> a_device(a.data(), a_size);
        sycl::buffer<int> c_device(c.data(), a_size);
        sycl::image<2> b_device(b.data(), sycl::image_channel_order::rgba,
                                sycl::image_channel_type::fp32, sycl::range<2>(8, 8));
```

(continues on next page)

(continued from previous page)

```

d_queue.submit([&](sycl::handler &cgh) {
    sycl::accessor<int, 1, sycl::access::mode::discard_write,
        sycl::access::target::global_buffer> c_res(c_device, cgh);
    sycl::accessor<int, 1, sycl::access::mode::read,
        sycl::access::target::constant_buffer> a_res(a_device, cgh);
    sycl::accessor<sycl::float4, 2, sycl::access::mode::write,
        sycl::access::target::image> b_res(b_device, cgh);

    sycl::float4 init = {0.f, 0.f, 0.f, 0.f};

    cgh.parallel_for<class ex1>(a_size, [=](sycl::id<1> idx) {
        c_res[idx] = a_res[idx];
        b_res.write(sycl::int2(0,0), init);
    });

});

}
return 0;
}

```

- Lines 8 and 9 contain the host allocations of arrays a, b, & c. The declaration of b is as a float4 because it will be accessed as an image on the device side.
- Lines 27 and 28 create an accessor for c_device that has an access mode of discard_write and a target of global_buffer.
- Lines 29 and 30 create an accessor for a_device that has an access mode of read and a target of constant_buffer.
- Lines 31 and 32 create an accessor for b_device that has an access mode of write and a target of image.

The accessors specify where and how the kernel will access these memory objects. The runtime is responsible for placing the memory objects in the correct location. Therefore, the runtime may copy data between host and device to meet the semantics of the accessor target.

Designate accessor targets to optimize the locality of access for a particular algorithm. For example, specify that local memory should be employed if much of the kernel access would benefit from memory that resides closer to the processing elements.

If the kernel attempts to violate the communicated accessor by either attempting to write on a read accessor or read on a write accessor, a compiler diagnostic is emitted. Not all combinations of access targets and access modes are compatible. For details, see the SYCL Specification.

2.4.1 Memory Objects

Memory objects are either buffers or images.

- Buffer object - a one-, two-, or three-dimensional array of elements. Buffers can be accessed via lower level C++ pointer types. For further information on buffers, see the SYCL Specification.
- Image object - a formatted opaque memory object stored in a type specific and optimized fashion. Access occurs through built-in functions. Image objects typically pertain to pictures comprised of pixels stored in a format such as RGB (red, green, blue intensity). For further information on images, see the SYCL Specification.

2.4.2 Accessors

Accessors provide access to buffers and images in the host or inside the kernel and also communicate data dependencies between the application and different kernels. The accessor communicates the data type, the size, the target, and the access mode. To enable good performance, pay particular attention to the target because the accessor specifies the memory type from the choices in the SYCL memory model.

The targets associated with buffers are:

- `global_buffer`
- `host_buffer`
- `constant_buffer`
- `local`

The targets associated with images are:

- `image`
- `host_image`
- `image_array`

Image access must also specify a channel order to communicate the format of the data being read. For example, an image may be specified as a `float4`, but accessed with a channel order of `RGBA`.

The access mode impacts correctness as well as performance and is one of `read`, `write`, `read_write`, `discard_write`, `discard_read_write`, or `atomic`. Mismatches in access mode and actual memory operations such as a `write` to a buffer with access mode `read` can result in compiler diagnostics as well as erroneous program state. The `discard_write` and `discard_read_write` access modes can provide performance benefits for some implementations. For further details on accessors, see the SYCL Specification.

2.4.3 Synchronization

It is possible to access a `buffer` without employing an `accessor`, however it should be the rare case. To do so safely, a `mutex_class` should be passed when a `buffer` is instantiated. For further details on this method, see the SYCL Specification.

Access Targets

| Target | Description |
|-----------------|------------------------------------------------------------------------------------------|
| host_buffer | Access the buffer on the host. |
| global_buffer | Access the buffer through global memory on the device. |
| constant_buffer | Access the buffer from constant memory on the device. This may enable some optimization. |
| local | Access the buffer from local memory on the device. |
| image | Access the image |
| image_array | Access an array of images |
| host_image | Access the image on the host. |

Access Modes

| Memory Access Mode | Description |
|--------------------|------------------------------------------------|
| read | Read-only |
| write | Write-only |
| read_write | Read and write |
| discard_write | Write-only access. Previous value is discarded |
| discard_read_write | Read and write. Previous value is discarded |
| atomic | Provide atomic, one at a time, access. |

2.4.4 Unified Shared Memory

An extension to the standard SYCL memory model is unified shared memory, which enables the sharing of memory between the host and device without explicit accessors in the source code. Instead, manage access and enforces dependencies with explicit functions to wait on events or signaling a `depends_on` relationship between events.

Another characteristic of unified shared memory is that it provides a C++ pointer-based alternative to the buffer programming model. Unified shared memory provides both explicit and implicit models for managing memory. In the explicit model, programmers are responsible for specifying when data should be copied between memory allocations on the host and allocations on a device. In the implicit model, the underlying runtime and device drivers are responsible for automatically migrating memory between the host and a device. Since unified shared memory does not rely on accessors, dependencies between operations must be specified using events. Programmers may either explicitly wait on event objects or use the `depends_on` method inside a command group to specify a list of events that must complete before a task may begin.

2.4.5 Memory Types

| Memory Type | Description |
|-----------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Constant Memory | A region of global memory that remains constant during the execution of a kernel. The host allocates and initializes memory objects placed into constant memory. |
| Global Memory | Accessible to all work-items in all work-groups. Read/write, may be cached, persistent across kernel invocations. |
| Local Memory | Shared between work-items in a single work-group and inaccessible to work-items in other work-groups. Example: Shared local memory on Intel HD Graphics 530 |
| Private Memory | A region of memory private to a work-item. Variables defined in one work-item's private memory are not visible to another work-item. Example: Register File on Intel HD Graphics 530 |

2.5 Kernel Programming Model

The kernel programming model for oneAPI is based upon the SYCL* kernel programming model. It enables explicit parallelism between the host and device. The parallelism is explicit in the sense that the programmer determines what code executes on the host and device; it is not automatic. The kernel code executes on the accelerator. Programs employing the oneAPI programming model support single source, meaning the host code and device code can be in the same source file. However, there are differences between the source code accepted in the host code and the device code with respect to language conformance and language features. The SYCL Specification defines in detail the required language features for host code and device code. The following is a summary that is specific to the oneAPI product.

2.5.1 C++ Language Requirements

The host code can be compiled by C++11 and later compilers and take advantage of supported C++11 and later language features. The device code requires a compiler that accepts all C++03 language features and the following C++11 features:

- Lambda expressions
- Variadic templates
- Alias templates
- rvalue references
- `std::function`, `std::string`, `std::vector`

In addition, the device code cannot use the following features:

- Virtual Functions
- Virtual Inheritance
- Exceptions handling – throws and catches between host and device
- Run Time Type Information (RTTI)
- Object management employing new and delete operators

The device code is specified via one of three language constructs: lambda expression, functor, or kernel class. The separation of host code and device code via these language constructs is natural and accomplished without language extensions. These different forms of expressing kernels give the developer flexibility in enmeshing the host code and device code. For example:

- To put the kernel code in line with the host code, consider employing a lambda expression.
- To have the device code separate from the host code, but still maintain the single source property, consider employing a functor.
- To port code from OpenCL programs or to ensure a more rigid separation between host and device code, consider employing the kernel class.

The Device code inside a lambda expression, functor, or kernel object can then specify the amount of parallelism to request through several mechanisms.

- `single_task` – execute a single instance of the kernel with a single work item.
- `parallel_for` – execute a kernel in parallel across a range of processing elements. Typically, this version of `parallel_for` is employed on “embarrassingly parallel” workloads.
- `parallel_for_work_group` – execute a kernel in parallel across a hierarchical range of processing elements using local memory and barriers.

The following code sample shows two combinations of invoking kernels:

1. `single_task` and C++ lambda (lines 33-40)
2. `parallel_for` and functor (lines 8-20 and line 50)

These constructs enclose the aforementioned kernel scope. For details, see the SYCL Specification.

```
#include <vector>
#include <CL/sycl.hpp>

#define SIZE 1024

namespace sycl = cl::sycl;

template <typename T> class Vassign {
    T val;
    sycl::accessor<T, 1, sycl::access::mode::read_write,
        sycl::access::target::global_buffer> access;

public:
    Vassign(T val_, sycl::accessor<T, 1, sycl::access::mode::read_write,
        sycl::access::target::global_buffer> &access_) : val(val_),
        access(access_) {}

    void operator()(sycl::id<1> id) { access[id] = 1; }

};

int main() {
    std::array<int, SIZE> a;

    for (int i = 0; i<SIZE; ++i) {
        a[i] = i;
    }
    {
        sycl::range<1> a_size{SIZE};
        sycl::buffer<int, 1> a_device(a.data(), a_size);
        sycl::queue d_queue;

        d_queue.submit([&](sycl::handler &cgh) {
            auto a_in = a_device.get_access<sycl::access::mode::write>(cgh);

            cgh.single_task<class ex1>([=]() {
                a_in[0] = 2;
            });
        });
    }
}
```

(continues on next page)

(continued from previous page)

```

{
    sycl::range<1> a_size{SIZE};
    sycl::buffer<int, 1> a_device(a.data(), a_size);
    sycl::queue d_queue;
    d_queue.submit([&](sycl::handler &cgh) {
        auto a_in = a_device.get_access<sycl::access::mode::read_write,
            sycl::access::target::global_buffer>(cgh);
        Vassign<int> F(0, a_in);
        cgh.parallel_for(sycl::range<1>(SIZE), F);
    });
}
}

```

2.5.2 Error Handling

C++ exception handling is the basis for handling error conditions in the programming model. Some restrictions on exceptions are in place due to the asynchronous nature of host and device execution. For example, it is not possible to throw an exception in kernel scope and catch it (in the traditional sense) in application scope. Instead, there are a set of restrictions and expectations in place when performing error handling. These include:

- At application scope, the full C++ exception handling mechanisms and capability are valid as long as there is no expectation that exceptions can cross to kernel scope.
- At the command group scope, exceptions are asynchronous with respect to the application scope. During command group construction, an `async_handler` can be declared to handle any exceptions occurring during execution in the command group.

For further details on error handling, see the SYCL Specification.

2.5.3 Fall Back

Typically, a command group is submitted and executed on the designated command queue; however, there may be cases where the command queue is unable to execute the group. In these cases, it is possible to specify a fall back command queue for the command group to be executed upon. This capability is handled by the runtime. This fallback mechanism is detailed in the SYCL Specification.

The following code fails due to the size of the workgroup when executed on Intel Processor Graphics, such as Intel HD Graphics 530. The SYCL specification allows specifying a secondary queue as a parameter to the submit function and this secondary queue is used if the device kernel runs into issues with submission to the first device.

```

#include<CL/sycl.hpp>
#include<iostream>
#define N 1024
#define M 32
using namespace cl::sycl;
int main() {
    {
        cpu_selector cpuSelector;
        queue cpuQueue(cpuSelector);
        queue defaultQueue;
        buffer<int, 2> buf(range<2>(N, N));
        defaultQueue.submit([&](handler &cgh) {
            auto bufacc = buf.get_access<access::mode::read_write>(cgh);

```

(continues on next page)

(continued from previous page)

```
cgh.parallel_for<class ndim>(nd_range<2>(range<2>(N,N),
    range<2>(M,M)), [=](nd_item<2> i){
    id<2> ind = i.get_global_id();
    bufacc[ind[0]][ind[1]] = ind[0]+ind[1];
});
},cpuQueue);
auto bufaccl = buf.get_access<access::mode::read>();
for(int i = 0; i < N; i++){
    for(int j = 0; j < N; j++){
        if(bufaccl[i][j] != i+j){
            std::cout<<"Wrong result\n";
            return 1;
        }
    }
}
std::cout<<"Correct results\n";
return 0;
}
```


LANGUAGE

Todo: C++ version minimum

Programming Model and subsections documented the C++ language features accepted in code at application scope and command group scope in a DPC++ program. Application scope and command group scope includes the code that executes on the host. That section also documented the C++ language features accepted in code at kernel scope in a DPC++ program. Kernel scope is the code that executes on the device. In general, the full capabilities of C++ are available at application and command group scope. At kernel scope there are limitations in accepted C++ due to the more limited, but focused, capabilities of accelerators.

Compilers from different vendors have small eccentricities or differences in their conformance to the C++ standard. The Intel oneAPI DPC++ Compiler is a LLVM-based compiler and therefore drafts the specific behavior of the LLVM-based compilers in accepting and creating executables from C++ source code. To determine the specific LLVM version that the Intel oneAPI DPC++ Compiler is based upon, use the `--version` option.

```
dpcpp --version
```

For example:

```
DPC++ Compiler 2021.1 (2019.8.x.0.1010)
Target: x86_64-pc-windows-msvc
Thread model: posix
InstalledDir: c:\PROGRA~2\INTELO~1\compiler\latest\windows\bin
```

3.1 Keywords

One of the design goals of DPC++ and SYCL is to not add keywords to the language. The motivation is to enable easier compiler vendor adoption. Whereas OpenCL C code and other accelerator-targeted languages require proprietary keywords, DPC++ does not.

3.2 Preprocessor Directives and Macros

Standard C++ preprocessing directives and macros are supported by the compiler. In addition, the SYCL Specification defines the SYCL specific preprocessor directives and macros.

The following preprocessor directives and macros are supported by the compiler.

| Directive | Description |
|--------------------------------------------------------|--------------------------------------------------------------|
| <code>SYCL_DUMP_IMAGES</code> | If true, instructs runtime to dump the device image |
| <code>SYCL_USE_KERNEL_SPV=<device binary></code> | Employ device binary to fulfill kernel launch request |
| <code>SYCL_PROGRAM_BUILD_OPTIONS</code> | Used to pass additional options for device program building. |

3.3 Standard Library Classes Required for the Interface

Programming for oneAPI employs a variety of vectors, strings, functions, and pointer objects common in STL programming.

The SYCL specification documents a facility to enable vendors to provide custom optimized implementations. Implementations require aliases for several STL interfaces. These are summarized as follows:

```
template<class T, class Alloc = std::allocator<T>>
using vector_class = std::vector<T, Alloc>
using string_class = std::string
template<class R, class ...ArgTypes>
using function_class = std::function<R (ArgTypes...)>
using mutex_class = std::mutex
template<class T>
using shared_ptr_class = std::shared_ptr<T>
template<class T>
using unique_ptr_class = std::unique_ptr<T>
template<class T>
using weak_ptr_class = std::weak_ptr<T>
template<class T>
using hash_class = std::hash<T>
using exception_ptr_class = std::exception_ptr
```

Template Parameters

- `T` –
- `allocatorT` –

PROGRAMMING INTERFACE

The Data Parallel C++ (DPC++) programming language and runtime consists of a set of C++ classes, templates, and libraries used to express a DPC++ program. This chapter provides a summary of the key classes, templates, and runtime libraries used to program.

4.1 Header File

4.2 Namespaces

4.3 Class Descriptions

The following sections provide further details on these items. These sections do not provide the exhaustive details found in the SYCL Specification. Instead, these sections provide:

- A summary that includes a description and the purpose
- Comments on the different constructors, if applicable
- Member function information, if applicable
- Special cases to consider with the DPC++ implementation compared to the SYCL Specification

For further details on SYCL, see the [SYCL Specification](#).

4.3.1 Runtime Classes

4.3.1.1 `device_selector`

```
class device_selector();
```

Member functions

| | |
|----------------------|--|
| <i>(constructor)</i> | |
| <i>(destructor)</i> | |
| <i>select_device</i> | |

Non-member functions

| | |
|-------------------------|--|
| <code>operator()</code> | |
|-------------------------|--|

(constructor)

```
device_selector(const device_selector &rhs);  
device_selector &operator=(const device_selector &rhs);
```

(destructor)

```
virtual ~device_selector();
```

select_device

```
device select_device() const;
```

Return value

operator()

```
virtual int operator()(const device &device) const = 0;
```

Return value

4.3.1.2 platform

```
class platform;
```

Abstraction for SYCL platform.

Member functions

| | |
|----------------------|-------------------------------------------|
| <i>(constructor)</i> | constructs a platform |
| <i>(destructor)</i> | destroys a platform |
| <i>get</i> | returns OpenCL platform ID |
| <i>get_devices</i> | returns devices bound to the platform |
| <i>get_info</i> | queries properties |
| <i>has_extension</i> | checks if platform has an extension |
| <i>is_host</i> | checks if platform has a SYCL host device |

Non-member functions

| | |
|----------------------|---------------------------------------|
| <i>get_platforms</i> | returns platforms bound to the system |
|----------------------|---------------------------------------|

Example

Demonstrates several methods for platform

```
#include <CL/sycl.hpp>

namespace sycl = cl::sycl;

int main() {
    auto platforms = sycl::platform::get_platforms();

    for (auto &platform : platforms) {
        std::cout << "Platform: "
                    << platform.get_info<sycl::info::platform::name>()
                    << std::endl;

        auto devices = platform.get_devices();
        for (auto &device : devices) {
            std::cout << "  Device: "
                        << device.get_info<sycl::info::device::name>()
                        << std::endl;
        }
    }

    return 0;
}
```

Output:

```
Platform: Intel(R) FPGA Emulation Platform for OpenCL(TM)
  Device: Intel(R) FPGA Emulation Device
Platform: Intel(R) OpenCL
  Device: Intel(R) Core(TM) i5-7300U CPU @ 2.60GHz
Platform: Intel(R) CPU Runtime for OpenCL(TM) Applications
  Device: Intel(R) Core(TM) i5-7300U CPU @ 2.60GHz
Platform: SYCL host platform
  Device: SYCL host device
```

(constructor)

```
platform();  
explicit platform(cl_platform_id platformID); 1  
explicit platform(const device_selector &deviceSelector); 2
```

Constructs a platform handle.

Arguments

| | |
|----------------|-------------------------------------------|
| platformID | OpenCL platform ID |
| deviceSelector | Platform must contain the selected device |

get

```
cl_platform_id get() const;
```

Returns OpenCL platform id used in the constructor.

get_devices

```
vector_class<device> get_devices(  
    info::device_type = info::device_type::all) const;
```

Returns vector of devices of the requested type

Arguments

| | |
|-------------|--------------------------------|
| device_type | limits type of device returned |
|-------------|--------------------------------|

Return value

vector containing devices of the specified type bound to the platform.

Example

See *Example*.

¹ Constructs a SYCL platform that retains an OpenCL id

² Selects a platform that contains the desired device

get_info

```
template< info::platform param >
typename info::param_traits<info::platform, param>::return_type get_info() const;
```

Returns information about the platform, as specified by `param`.

Return value

Requested information

Example

See *Example*.

has_extension

```
bool has_extension(const string_class &extension) const;
```

Checks if the platform has the requested extension.

Arguments

| | |
|-----------|--|
| extension | |
|-----------|--|

Return value

true if the platform has extension

is_host

```
bool is_host() const;
```

Checks if the platform contains a SYCL *host device*

Return value

true if the platform contains a host device

get_platforms

```
static vector_class<platform> get_platforms();
```

Returns vector of platforms

Return value

vector_class containing SYCL platforms bound to the system

Example

See *Example*.

4.3.1.3 context

```
class context;
```

Member functions

| | |
|----------------------|---------------------------------------|
| <i>(constructor)</i> | constructs a context |
| <i>get</i> | returns OpenCL context ID |
| <i>is_host</i> | checks if contains a SYCL host device |
| <i>get_platform</i> | |
| <i>get_devices</i> | returns devices bound to the context |
| <i>get_info</i> | queries properties |

(constructor)

```
explicit context(const property_list &propList = {});
context(async_handler asyncHandler,
        const property_list &propList = {});
context(const device &dev, const property_list &propList = {});
context(const device &dev, async_handler asyncHandler,
        const property_list &propList = {});
context(const platform &plt, const property_list &propList = {});
context(const platform &plt, async_handler asyncHandler,
        const property_list &propList = {});
context(const vector_class<device> &deviceList,
        const property_list &propList = {});
context(const vector_class<device> &deviceList,
        async_handler asyncHandler, const property_list &propList = {});
context(cl_context clContext, async_handler asyncHandler = {});
```


Arguments

| | |
|--------------|--|
| propList | |
| asyncHandler | |
| dev | |
| plt | |
| deviceList | |

get

```
cl_context get() const;
```

Return value

is_host

```
bool is_host() const;
```

Return value

get_platform

```
platform get_platform() const;
```

Return value

get_devices

```
vector_class<device> get_devices() const;
```

Return value

get_info

```
template <info::context param>
typename info::param_traits<info::context, param>::return_type get_info() const;
```

Return value

4.3.1.4 device

```
class device;
```

Member functions

| | |
|---------------------------|--|
| <i>(constructor)</i> | |
| <i>(destructor)</i> | |
| <i>get</i> | |
| <i>is_host</i> | |
| <i>is_cpu</i> | |
| <i>is_gpu</i> | |
| <i>is_accelerator</i> | |
| <i>get_platform</i> | |
| <i>get_info</i> | |
| <i>has_extension</i> | |
| <i>create_sub_devices</i> | |

Non-member functions

| | |
|--------------------|--|
| <i>get_devices</i> | |
|--------------------|--|

(constructor)

```
device();  
explicit device(cl_device_id deviceId);  
explicit device(const device_selector &deviceSelector);
```

Arguments

| | |
|----------------|--|
| deviceId | |
| deviceSelector | |

get

```
cl_device_id get() const;
```

Return value**is_host**

```
bool is_host() const;
```

Return value**is_cpu**

```
bool is_cpu() const;
```

Return value**is_gpu**

```
bool is_gpu() const;
```

Return value**is_accelerator**

```
bool is_accelerator() const;
```

Return value**get_platform**

```
platform get_platform() const;
```

Return value**get_info**

```
template <info::device param>  
typename info::param_traits<info::device, param>::return_type  
get_info() const;
```

Return value

Example

See *Example*.

has_extension

```
bool has_extension(const string_class &extension) const;
```

Arguments

| | |
|-----------|--|
| extension | |
|-----------|--|

Return value

create_sub_devices

```
template <info::partition_property prop>  
vector_class<device> create_sub_devices(size_t nbSubDev) const; 4
```

```
template <info::partition_property prop>  
vector_class<device> create_sub_devices(const vector_class<size_t> &counts) └  
    ↪ const; 5
```

```
template <info::partition_property prop>  
vector_class<device> create_sub_devices(info::affinity_domain affinityDomain) └  
    ↪ const; 6
```

Arguments

| | |
|----------------|--|
| nbSubDev | |
| counts | |
| affinityDomain | |

⁴ Available only when `prop == info::partition_property::partition_equally`

⁵ Available only when `prop == info::partition_property::partition_by_counts`

⁶ Available only when `prop == info::partition_property::partition_by_affinity_domain`

Return value

get_devices

```
static vector_class<device> get_devices(
    info::device_type deviceType = info::device_type::all);
```

Return value

4.3.1.5 queue

```
class queue;
```

Member functions

| | |
|---------------------------|--|
| <i>(constructor)</i> | |
| (destructor) | |
| <i>get</i> | |
| <i>get_context</i> | |
| <i>get_device</i> | |
| <i>get_info</i> | |
| <i>is_host</i> | |
| <i>submit</i> | |
| <i>wait</i> | |
| <i>wait_and_throw</i> | |
| <i>throw_asynchronous</i> | |

(constructor)

```
explicit queue(const property_list &propList = {});
explicit queue(const async_handler &asyncHandler,
               const property_list &propList = {});
explicit queue(const device_selector &deviceSelector,
               const property_list &propList = {});
explicit queue(const device_selector &deviceSelector,
               const async_handler &asyncHandler,
               const property_list &propList = {});
explicit queue(const device &syclDevice, const property_list &propList = {});
explicit queue(const device &syclDevice, const async_handler &asyncHandler,
               const property_list &propList = {});
explicit queue(const context &syclContext,
               const device_selector &deviceSelector,
               const property_list &propList = {});
explicit queue(const context &syclContext,
               const device_selector &deviceSelector,
               const async_handler &asyncHandler,
               const property_list &propList = {});
explicit queue(const context &syclContext,
               const device &syclDevice,
```

(continues on next page)

(continued from previous page)

```
        const property_list &propList = {});  
explicit queue(const context &syclContext, const device &syclDevice,  
               const async_handler &asyncHandler,  
               const property_list &propList = {});  
explicit queue(cl_command_queue clQueue, const context& syclContext,  
               const async_handler &asyncHandler = {});
```

get

```
cl_command_queue get() const;
```

Return value**get_context**

```
context get_context() const;
```

Return value**get_device**

```
device get_device() const;
```

Return value**is_host**

```
bool is_host() const;
```

Return value**get_info**

```
template <info::queue param>  
typename info::param_traits<info::queue, param>::return_type get_info() const;
```

Return value**submit**

```
template <typename T>
event submit(T cgf);

template <typename T>
event submit(T cgf, const queue &secondaryQueue);
```

Arguments

| | |
|----------------|--|
| cgf | |
| secondaryQueue | |

Return value**wait**

```
void wait();
```

wait_and_throw

```
void wait_and_throw();
```

throw_asynchronous

```
void throw_asynchronous();
```

4.3.1.6 event

```
class event;
```

Member functions

| | |
|---------------------------|--|
| <i>(constructor)</i> | |
| <i>(destructor)</i> | |
| <i>cl_event_get</i> | |
| <i>is_host</i> | |
| <i>get_wait_list</i> | |
| <i>wait</i> | |
| <i>wait_and_throw</i> | |
| <i>get_info</i> | |
| <i>get_profiling_info</i> | |

(constructor)

```
event();  
event(cl_event clEvent, const context& syclContext);
```

cl_event_get

```
cl_event get();
```

Return value

is_host

```
bool is_host() const;
```

Return value

get_wait_list

```
vector_class<event> get_wait_list();
```

Return value

wait

```
void wait();  
static void wait(const vector_class<event> &eventList);
```

wait_and_throw

```
void wait_and_throw();  
static void wait_and_throw(const vector_class<event> &eventList);
```


get_info

```
template <info::event param>
typename info::param_traits<info::event, param>::return_type get_info() const;
```

Return value**get_profiling_info**

```
template <info::event_profiling param>
typename info::param_traits<info::event_profiling, param>::return_type get_profiling_
↪info() const;
```

Return value**4.3.2 Data Access****4.3.2.1 accessor**

```
template<
    typename dataT,
    int dimensions,
    access::mode accessmode,
    access::target accessTarget = access::target::global_buffer,
    access::placeholder isPlaceholder = access::placeholder::false_t
> accessor;
```

A DPC++ `accessor` encapsulates reading and writing memory objects which can be buffers, images, or device local memory. Creating an `accessor` requires a method to reference the desired access target. Construction also requires the type of the memory object, the dimensionality of the memory object, the access mode, and a placeholder argument.

Template parameters

| | |
|----------------------------|--------------------------|
| <code>dataT</code> | type of buffer element |
| <code>dimensions</code> | dimensionality of buffer |
| <code>accessmode</code> | type of access |
| <code>accessTarget</code> | type of memory |
| <code>isPlaceholder</code> | placeholder |

Member types

| | |
|-----------------|--------------|
| value_type | dataT |
| reference | dataT& |
| const_reference | const dataT& |

Member functions

| | |
|------------------|------------------------|
| (constructor) | constructs an accessor |
| (destructor) | destroys the accessor |
| is_placeholder | |
| <i>get_size</i> | |
| <i>get_count</i> | |
| get_range | |
| get_offset | |

get_size

```
size_t get_size() const
```

Description

get_count

```
size_t get_size() const 1  
size_t get_size(int b) const 2  
size_t get_size(int c, int d) const 3
```

Description

A common method of construction can employ the `get_access` method of the memory object to specify the object and infer the other parameters from that memory object.

See [Accessors](#) for access modes and targets.

Placeholder accessors are those created independent of a command group and then later associated with a particular memory object. Designation of a placeholder accessor is communicated via the placeholder argument set to `access::placeholder::true_t` if so and `access::placeholder::false_t` otherwise.

Once an accessor is created, query member functions to review accessor information. These member functions include:

- `is_placeholder` – return true if accessor is a placeholder, not yet associated with a memory object, false otherwise
- `get_size` – obtain the size (in bytes) of the memory object
- `get_count` – obtain the number of elements of the memory object
- `get_range` – obtain the range of the memory object, where range is a range class
- `get_offset` – obtain the offset of the memory object

¹ No arguments

² single argument

³ 2 arguments

An accessor can reference a subset of a memory object; this is the offset of the accessor into the memory object.

4.3.2.2 Atomic

The DPC++ `atomic` class encapsulates operations and member functions to guarantee synchronized access to data values. Construction of an atomic object requires a reference to a *Address Spaces*. A `multi_ptr` is an abstraction on top of a low-level pointer that enables efficient access across the host and devices.

The atomic member functions are modeled after the C++ standard atomic functions. They are documented more fully in the SYCL Specification and include the following:

- `Store` – store a value
- `Load` – load a value
- `Exchange` – swap two values
- `Compare_exchange_strong` - compares two values for equality and exchanges based on result
- `Fetch_add` – add a value to the value pointed to by a `multi_ptr`
- `Fetch_sub` - subtract a value from the value pointed to by a `multi_ptr`
- `Fetch_and` – bitwise and a value from the value pointed to by a `multi_ptr`
- `Fetch_or` - bitwise or a value from the value pointed to by a `multi_ptr`
- `Fetch_xor` - bitwise xor a value from the value pointed to by a `multi_ptr`
- `Fetch_min` – compute the minimum between a value and the value pointed to by a `multi_ptr`
- `Fetch_max` - compute the maximum between a value and the value pointed to by a `multi_ptr`

In addition to the member functions above, a set of functions with the same capabilities are available acting on atomic types. These functions are similarly named with the addition of “`atomic_`” prepended.

```
enum class memory_order
```

```
    enumerator relaxed
```

```
template<typename T, access::address_space addressSpace = access::address_space::global_space>
class atomic
```

```
    tparam T
```

```
    tparam addressSpace
```

```
template<typename pointerT>
```

```
    atomic (multi_ptr<pointerT, addressSpace> ptr)
```

```
    Template Parameters pointerT –
```

```
    Parameters ptr –
```

```
void store (T operand, memory_order memoryOrder = memory_order::relaxed)
```

```
T load (memory_order memoryOrder = memory_order::relaxed) const
```

```
T exchange (T operand, memory_order memoryOrder = memory_order::relaxed)
```

```
    Parameters
```

```
    • operand –
```

```
    • memoryOrder –
```

```
bool compare_exchange_strong (T &expected, T desired, memory_order successMemoryOrder = memory_order::relaxed, memory_order failMemoryOrder = memory_order::relaxed)
```

Parameters

- **expected** –
- **desired** –
- **successMemoryOrder** –
- **failMemoryOrder** –

Available only when: $T \neq \text{float}$

```
T fetch_add (T operand, memory_order memoryOrder = memory_order::relaxed)
T fetch_sub (T operand, memory_order memoryOrder = memory_order::relaxed)
T fetch_and (T operand, memory_order memoryOrder = memory_order::relaxed)
T fetch_or (T operand, memory_order memoryOrder = memory_order::relaxed)
T fetch_xor (T operand, memory_order memoryOrder = memory_order::relaxed)
T fetch_min (T operand, memory_order memoryOrder = memory_order::relaxed)
T fetch_max (T operand, memory_order memoryOrder = memory_order::relaxed)
```

Parameters

- **operand** –
- **memoryOrder** –

Available only when: $T \neq \text{float}$

4.3.2.3 Buffer

A DPC++ `buffer` encapsulates a 1-, 2-, or 3-dimensional array that is shared between host and devices. Creating a buffer requires the number of dimensions of the array as well as the type of the underlying data.

The class contains multiple constructors with different combinations of ranges, allocators, and property lists.

- The memory employed by the buffer is already existing in host memory. In this case, a pointer to the memory is passed to the constructor.
- Temporary memory is allocated for the buffer by employing the constructors that do not include a `hostData` parameter.
- An allocator object is passed, which provides an alternative memory allocator to be used for allocating the temporary memory for the buffer. Special arguments, termed properties, can be provided to the constructor for cases where host memory use is desired (`use_host_ptr`), use of the `mutex_class` is desired (`use_mutex`), and single context only (`context_bound`) is desired.

Once a buffer is allocated, query member functions to learn more. These member functions include:

- `get_range` – obtain the range object associated with the buffer
- `get_count` – obtain the number of elements in the buffer
- `get_size` – obtain the size of the buffer in bytes
- `get_allocator` – obtain the allocator that was provided in creating the buffer
- `is_sub_buffer` – return if buffer is a sub-buffer or not

```
template<typename T, int dimensions = 1, typename AllocatorT = cl::sycl::buffer_allocator>
```

class buffer**Template Parameters**

- **T** – buffer element type
- **dimensions** – dimensionality of the data
- **AllocatorT** – STL allocator

```
using value_type = T
```

```
using reference = value_type&
```

```
using const_reference = const value_type&
```

```
using allocator_type = AllocatorT
```

Standard types

```
buffer (const range<dimensions> &bufferRange, const property_list &propList = {})
```

```
buffer (const range<dimensions> &bufferRange, AllocatorT allocator, const property_list &propList = {})
```

```
buffer (T *hostData, const range<dimensions> &bufferRange, const property_list &propList = {})
```

```
buffer (T *hostData, const range<dimensions> &bufferRange, AllocatorT allocator, const property_list &propList = {})
```

```
buffer (const T *hostData, const range<dimensions> &bufferRange, const property_list &propList = {})
```

```
buffer (const T *hostData, const range<dimensions> &bufferRange, AllocatorT allocator, const property_list &propList = {})
```

```
buffer (const shared_ptr_class<T> &hostData, const range<dimensions> &bufferRange, AllocatorT allocator, const property_list &propList = {})
```

```
buffer (const shared_ptr_class<T> &hostData, const range<dimensions> &bufferRange, const property_list &propList = {})
```

```
template<class InputIterator>
```

```
buffer<T, 1> (InputIterator first, InputIterator last, AllocatorT allocator, const property_list &propList = {})
```

```
template<class InputIterator>
```

```
buffer<T, 1> (InputIterator first, InputIterator last, const property_list &propList = {})
```

```
buffer (buffer<T, dimensions, AllocatorT> b, const id<dimensions> &baseIndex, const range<dimensions> &subRange)
```

Constructors.

Parameters

- **bufferRange** – Dimensions of buffer
- **propList** – Properties for buffer
- **hostData** – Home for data
- **allocator** – STL allocator for data
- **first** – first and last are STL iterators used to initialize the buffer
- **last** – first and last are STL iterators used to initialize the buffer
- **b** – Use buffer *b* to initialize
- **baseIndex** – Index in *b*
- **subRange** – Range of data in *b*

```
buffer (cl_mem clMemObject, const context &syclContext, event availableEvent = {})
```

Parameters

- `clMemObject` –
- `syclContext` –
- `availableEvent` –

Simplified buffer constructor when `dimensions == 1`.

```
range<dimensions> get_range() const
```

Returns

```
size_t get_count() const
```

Returns

```
size_t get_size() const
```

Returns

```
AllocatorT get_allocator() const
```

Returns

```
template<access::mode mode, access::target target = access::target::global_buffer>  
accessor<T, dimensions, mode, target> get_access(handler &commandGroupHandler)  
  
template<access::mode mode>  
accessor<T, dimensions, mode, access::target::host_buffer> get_access()  
  
template<access::mode mode, access::target target = access::target::global_buffer>  
accessor<T, dimensions, mode, target> get_access(handler &commandGroupHandler,  
                                                range<dimensions> accessRange,  
                                                id<dimensions> accessOffset = {})  
  
template<access::mode mode>  
accessor<T, dimensions, mode, access::target::host_buffer> get_access(range<dimensions> access-  
                                                                    Range, id<dimensions> ac-  
                                                                    cessOffset = {})
```

Template Parameters

- `mode` –
- `target` –

Parameters

- `handler` –
- `accessRange` –
- `accessOffset` –

Returns

```
template<typename Destination = std::nullptr_t>  
void set_final_data(Destination finalData = nullptr)
```

Template Parameters `Destination` –**Parameters** `finalData` –

```
void set_write_back(bool flag = true)
```

Parameters `flag` –

```
bool is_sub_buffer() const
```

Returns

```
template<typename ReinterpretT, int ReinterpretDim>
buffer<ReinterpretT, ReinterpretDim, AllocatorT> reinterpret (range<ReinterpretDim> reinterpretRange) const
```

Template Parameters

- **ReinterpretT** –
- **ReinterpretDim** –

Parameters **reinterpretRange** –

Returns

```
class use_host_ptr
```

```
use_host_ptr ()
```

```
class use_mutex
```

```
use_mutex (mutex_class &mutexRef)
```

Parameters **mutexRef** –

```
mutex_class *get_mutex_ptr () const
```

Returns

```
class context_bound
```

```
context_bound (context boundContext)
```

Parameters **boundContext** –

```
context get_context () const
```

Returns**4.3.2.4 Image**

A DPC++ image encapsulates a 1-, 2-, or 3-dimensional set of data shared between host and devices. Creating an image requires the number of dimensions of the array as well as the order and type of the underlying data.

The class contains multiple constructors with different combinations of orders, types, ranges, allocators, and property lists.

- The memory employed by the image is already existing host memory. In this case, a pointer to the memory is passed to the constructor.
- Temporary memory is allocated for the image by employing the constructors that do not include a *hostPointer* parameter.
- An allocator object is passed, which provides an alternative memory allocator to be used for allocating the temporary memory for the buffer.
- When host memory use is desired (*use_host_ptr*), use of the *mutex_class* is desired (*use_mutex*), and if a single context only (*context_bound*) is desired, special arguments, termed properties, are provided to the constructor.

Once a buffer is allocated, query member functions to learn more. These member functions include

- `get_range` – obtain the range object associated with the image
- `get_pitch` – obtain the range associated with a one-dimensional image
- `get_count` – obtain the number of elements in the image
- `get_size` – obtain the size of the image (in bytes)
- `get_allocator` – obtain the allocator that was provided in creating the image
- `get_access` – obtain an accessor to the image with the specified access mode and target

4.3.2.5 Address Spaces

The DPC++ `multi_pointer` class encapsulates lower level pointers that point to abstract device memory.

Constructors for the `multi_pointer` class enable explicit mention of the address space of the memory. The following lists the address space with the appropriate identifier:

- Global memory – `global_space`
- Local memory – `local_space`
- Constant memory – `constant_space`
- Private memory – `private_space`

The constructors can also be called in an unqualified fashion for cases where the location will be known later.

Member functions include standard pointer operations such as ++ (increment), – (decrement), + (plus), and – (minus). A prefetch function is also specified to aid in optimization and is implementation defined.

Conversion operations are also available to convert between the raw underlying pointer and an OpenCL program's C pointer for interoperability. Consult the SYCL Specification for complete details.

enum class address_space

```
enumerator global_space
enumerator local_space
enumerator constant_space
enumerator private_space
```

```
template<typename ElementType, access::address_space Space>
class multi_ptr
```

```
using element_type = ElementType
using difference_type = std::ptrdiff_t
using pointer_t = ElementType*
using const_pointer_t = const ElementType*
using reference_t = ElementType&
using const_reference_t = const ElementType&
```

Implementation defined pointer and reference types that correspond to SYCL/OpenCL interoperability types for OpenCL C functions

```
static constexpr access::address_space address_space = Space

multi_ptr()
multi_ptr(const multi_ptr&)
```



```

multi_ptr (multi_ptr&&)
multi_ptr (pointer_t)
multi_ptr (ElementType*)
multi_ptr (std::nullptr_t)
template<int dimensions, access::mode Mode, access::placeholder isPlaceholder>
multi_ptr (accessor<ElementType, dimensions, Mode, access::target::global_buffer, isPlaceholder>)
template<int dimensions, access::mode Mode, access::placeholder isPlaceholder>
multi_ptr (accessor<ElementType, dimensions, Mode, access::target::local, isPlaceholder>)
template<int dimensions, access::mode Mode, access::placeholder isPlaceholder>
multi_ptr (accessor<ElementType, dimensions, Mode, access::target::constant_buffer, isPlaceholder>)
~multi_ptr ()

```

Assignment and access operators

```

multi_ptr &operator= (const multi_ptr&)
multi_ptr &operator= (multi_ptr&&)
multi_ptr &operator= (pointer_t)
multi_ptr &operator= (ElementType*)
multi_ptr &operator= (std::nullptr_t)
friend ElementType &operator* (const multi_ptr &mp)
ElementType *operator-> () const
pointer_t get () const

```

Returns underlying OpenCL C pointer

```

operator ElementType* () const
    Implicit conversion to the underlying pointer type
operator multi_ptr<void, Space> () const
operator multi_ptr<const void, Space> () const
operator multi_ptr<const ElementType, Space> () const
    Implicit conversion to a multi_ptr
void prefetch (size_t numElements) const

```

Arithmetic operators

```

friend multi_ptr &operator++ (multi_ptr &mp)
friend multi_ptr operator++ (multi_ptr &mp, int)
friend multi_ptr &operator-- (multi_ptr &mp)
friend multi_ptr operator-- (multi_ptr &mp, int)
friend multi_ptr &operator+= (multi_ptr &lhs, difference_type r)
friend multi_ptr &operator-= (multi_ptr &lhs, difference_type r)
friend multi_ptr operator+ (const multi_ptr &lhs, difference_type r)
friend multi_ptr operator- (const multi_ptr &lhs, difference_type r)

```

Relational operators

```

friend bool operator==(const multi_ptr &lhs, const multi_ptr &rhs)
friend bool operator!=(const multi_ptr &lhs, const multi_ptr &rhs)
friend bool operator<(const multi_ptr &lhs, const multi_ptr &rhs)
friend bool operator>(const multi_ptr &lhs, const multi_ptr &rhs)
friend bool operator<=(const multi_ptr &lhs, const multi_ptr &rhs)
friend bool operator>=(const multi_ptr &lhs, const multi_ptr &rhs)
friend bool operator==(const multi_ptr &lhs, std::nullptr_t)
friend bool operator!=(const multi_ptr &lhs, std::nullptr_t)
friend bool operator<(const multi_ptr &lhs, std::nullptr_t)
friend bool operator>(const multi_ptr &lhs, std::nullptr_t)
friend bool operator<=(const multi_ptr &lhs, std::nullptr_t)
friend bool operator>=(const multi_ptr &lhs, std::nullptr_t)
friend bool operator==(std::nullptr_t, const multi_ptr &rhs)
friend bool operator!=(std::nullptr_t, const multi_ptr &rhs)
friend bool operator<(std::nullptr_t, const multi_ptr &rhs)
friend bool operator>(std::nullptr_t, const multi_ptr &rhs)
friend bool operator<=(std::nullptr_t, const multi_ptr &rhs)
friend bool operator>=(std::nullptr_t, const multi_ptr &rhs)

```

```

template<typename ElementType>
using global_ptr = multi_ptr<ElementType, access::address_space::global_space>
template<typename ElementType>
using local_ptr = multi_ptr<ElementType, access::address_space::local_space>
template<typename ElementType>
using constant_ptr = multi_ptr<ElementType, access::address_space::constant_space>
template<typename ElementType>
using private_ptr = multi_ptr<ElementType, access::address_space::private_space>

```

4.3.2.6 Samplers

4.3.3 Expressing Parallelism

4.3.3.1 Range

The DPC++ `range` class encapsulates the iteration domain of a work-group or the entire kernel dispatch. Constructors for a range object take one, two, or three arguments of `size_t` dependent on the dimensionality of the range, either one, two, or three dimensions respectively.

Member functions include:

- `get` – obtain the specified dimension
- `size` – obtain the size of the range

Additional functions allow construction of new ranges from old ranges with additional operations on the range. For example:

```
Range<2> +() const??
```

```
template<int dimensions = 1>
class range
```

```
    range (size_t dim0)
    range (size_t dim0, size_t dim1)
    range (size_t dim0, size_t dim1, size_t dim2)
```

Parameters

- **dim0** –
- **dim1** –
- **dim2** –

```
size_t get (int dimension) const
```

Parameters **dimension** –

Returns

```
size_t &operator[] (int dimension)
```

```
size_t operator[] (int dimension) const
```

Parameters **dimension** –

Returns

```
size_t size () const
```

```
friend range operator+ (const range &lhs, const range &rhs)
friend range operator- (const range &lhs, const range &rhs)
friend range operator* (const range &lhs, const range &rhs)
friend range operator/ (const range &lhs, const range &rhs)
friend range operator% (const range &lhs, const range &rhs)
friend range operator<< (const range &lhs, const range &rhs)
friend range operator>> (const range &lhs, const range &rhs)
friend range operator& (const range &lhs, const range &rhs)
friend range operator| (const range &lhs, const range &rhs)
friend range operator^ (const range &lhs, const range &rhs)
friend range operator&& (const range &lhs, const range &rhs)
friend range operator|| (const range &lhs, const range &rhs)
friend range operator< (const range &lhs, const range &rhs)
friend range operator> (const range &lhs, const range &rhs)
friend range operator<= (const range &lhs, const range &rhs)
friend range operator>= (const range &lhs, const range &rhs)

friend range operator+ (const range &lhs, const size_t &rhs)
friend range operator- (const range &lhs, const size_t &rhs)
friend range operator* (const range &lhs, const size_t &rhs)
friend range operator/ (const range &lhs, const size_t &rhs)
friend range operator% (const range &lhs, const size_t &rhs)
friend range operator<< (const range &lhs, const size_t &rhs)
friend range operator>> (const range &lhs, const size_t &rhs)
```

```
friend range operator& (const range &lhs, const size_t &rhs)
friend range operator| (const range &lhs, const size_t &rhs)
friend range operator^ (const range &lhs, const size_t &rhs)
friend range operator&& (const range &lhs, const size_t &rhs)
friend range operator|| (const range &lhs, const size_t &rhs)
friend range operator< (const range &lhs, const size_t &rhs)
friend range operator> (const range &lhs, const size_t &rhs)
friend range operator<= (const range &lhs, const size_t &rhs)
friend range operator>= (const range &lhs, const size_t &rhs)

friend range operator+ (const size_t &lhs, const range &rhs)
friend range operator- (const size_t &lhs, const range &rhs)
friend range operator* (const size_t &lhs, const range &rhs)
friend range operator/ (const size_t &lhs, const range &rhs)
friend range operator% (const size_t &lhs, const range &rhs)
friend range operator<< (const size_t &lhs, const range &rhs)
friend range operator>> (const size_t &lhs, const range &rhs)
friend range operator& (const size_t &lhs, const range &rhs)
friend range operator| (const size_t &lhs, const range &rhs)
friend range operator^ (const size_t &lhs, const range &rhs)
friend range operator&& (const size_t &lhs, const range &rhs)
friend range operator|| (const size_t &lhs, const range &rhs)
friend range operator< (const size_t &lhs, const range &rhs)
friend range operator> (const size_t &lhs, const range &rhs)
friend range operator<= (const size_t &lhs, const range &rhs)
friend range operator>= (const size_t &lhs, const range &rhs)

friend range operator+= (const range &lhs, const range &rhs)
friend range operator-= (const range &lhs, const range &rhs)
friend range operator*= (const range &lhs, const range &rhs)
friend range operator/= (const range &lhs, const range &rhs)
friend range operator%= (const range &lhs, const range &rhs)
friend range operator<<= (const range &lhs, const range &rhs)
friend range operator>>= (const range &lhs, const range &rhs)
friend range operator&= (const range &lhs, const range &rhs)
friend range operator|= (const range &lhs, const range &rhs)
friend range operator^= (const range &lhs, const range &rhs)

friend range operator+= (const range &lhs, const size_t &rhs)
friend range operator-= (const range &lhs, const size_t &rhs)
friend range operator*= (const range &lhs, const size_t &rhs)
friend range operator/= (const range &lhs, const size_t &rhs)
friend range operator%= (const range &lhs, const size_t &rhs)
friend range operator<<= (const range &lhs, const size_t &rhs)
friend range operator>>= (const range &lhs, const size_t &rhs)
friend range operator&= (const range &lhs, const size_t &rhs)
friend range operator|= (const range &lhs, const size_t &rhs)
friend range operator^= (const range &lhs, const size_t &rhs)
```

4.3.3.2 Nd_range

The DPC++ `nd_range` class encapsulates the iteration domain of the work-groups and kernel dispatch. It is the entire iteration space of data that a kernel may operation upon. The constructor for an `nd_range` object take the global range, local range, and an optional offset.

Member functions include:

- `get_global_range` – obtain the global range
- `get_local_range` – obtain the local range
- `get_group_range` – obtain the number of groups in each dimension of the `nd_range`
- `get_offset` – obtain the offset

```
template<int dimensions = 1>
class nd_range
```

```
    nd_range (range<dimensions> globalSize, range<dimensions> localSize, id<dimensions> offset =
              id<dimensions>())
    range<dimensions> get_global_range () const
    range<dimensions> get_local_range () const
    range<dimensions> get_group_range () const
    id<dimensions> get_offset () const
```

4.3.3.3 ID

The `id` class encapsulates a vector of dimensions that identify an index into a global or local range. Constructors for the class take one to three integer arguments representing a one, two, and three dimension ID. Each integer argument specifies the size of the dimension. ID objects can also be constructed as a placeholder where the dimension is unspecified and set to zero by default. Construction can also be based upon the dimension of an existing range or item.

The class supports operations such as + (plus), - (minus), and many more. Consult the SYCL Specification for complete details.

```
template<int dimensions = 1>
class id
```

tparam dimensions

```
    id ()
    id (size_t dim0)
    id (size_t dim0, size_t dim1)
    id (size_t dim0, size_t dim1, size_t dim2)
    id (const range<dimensions> &range)
    id (const item<dimensions> &item)
```

Parameters

- `dim0` –
- `dim1` –
- `dim2` –
- `range` –

- **item –**

`size_t get (int dimension) const`

Returns

`size_t &operator [] (int dimension)`

Returns

`size_t operator [] (int dimension) const`

Returns

`friend id operator+ (const id &lhs, const id &rhs)`

`friend id operator- (const id &lhs, const id &rhs)`

`friend id operator* (const id &lhs, const id &rhs)`

`friend id operator/ (const id &lhs, const id &rhs)`

`friend id operator% (const id &lhs, const id &rhs)`

`friend id operator<< (const id &lhs, const id &rhs)`

`friend id operator>> (const id &lhs, const id &rhs)`

`friend id operator& (const id &lhs, const id &rhs)`

`friend id operator| (const id &lhs, const id &rhs)`

`friend id operator^ (const id &lhs, const id &rhs)`

`friend id operator&& (const id &lhs, const id &rhs)`

`friend id operator|| (const id &lhs, const id &rhs)`

`friend id operator< (const id &lhs, const id &rhs)`

`friend id operator> (const id &lhs, const id &rhs)`

`friend id operator<= (const id &lhs, const id &rhs)`

`friend id operator>= (const id &lhs, const id &rhs)`

`friend id operator+ (const id &lhs, const size_t &rhs)`

`friend id operator- (const id &lhs, const size_t &rhs)`

`friend id operator* (const id &lhs, const size_t &rhs)`

`friend id operator/ (const id &lhs, const size_t &rhs)`

`friend id operator% (const id &lhs, const size_t &rhs)`

`friend id operator<< (const id &lhs, const size_t &rhs)`

`friend id operator>> (const id &lhs, const size_t &rhs)`

`friend id operator& (const id &lhs, const size_t &rhs)`

`friend id operator| (const id &lhs, const size_t &rhs)`

`friend id operator^ (const id &lhs, const size_t &rhs)`

`friend id operator&& (const id &lhs, const size_t &rhs)`

`friend id operator|| (const id &lhs, const size_t &rhs)`

`friend id operator< (const id &lhs, const size_t &rhs)`

`friend id operator> (const id &lhs, const size_t &rhs)`

`friend id operator<= (const id &lhs, const size_t &rhs)`

`friend id operator>= (const id &lhs, const size_t &rhs)`

`friend id operator+ (const size_t &lhs, const id &rhs)`

`friend id operator- (const size_t &lhs, const id &rhs)`

`friend id operator* (const size_t &lhs, const id &rhs)`

```

friend id operator/ (const size_t &lhs, const id &rhs)
friend id operator% (const size_t &lhs, const id &rhs)
friend id operator<< (const size_t &lhs, const id &rhs)
friend id operator>> (const size_t &lhs, const id &rhs)
friend id operator& (const size_t &lhs, const id &rhs)
friend id operator| (const size_t &lhs, const id &rhs)
friend id operator^ (const size_t &lhs, const id &rhs)
friend id operator&& (const size_t &lhs, const id &rhs)
friend id operator|| (const size_t &lhs, const id &rhs)
friend id operator< (const size_t &lhs, const id &rhs)
friend id operator> (const size_t &lhs, const id &rhs)
friend id operator<= (const size_t &lhs, const id &rhs)
friend id operator>= (const size_t &lhs, const id &rhs)

friend id operator+= (id &lhs, const id &rhs)
friend id operator*= (id &lhs, const id &rhs)
friend id operator/= (id &lhs, const id &rhs)
friend id operator%= (id &lhs, const id &rhs)
friend id operator<<= (id &lhs, const id &rhs)
friend id operator>>= (id &lhs, const id &rhs)
friend id operator&= (id &lhs, const id &rhs)
friend id operator|= (id &lhs, const id &rhs)
friend id operator^= (id &lhs, const id &rhs)

friend id operator+= (id &lhs, const size_t &rhs)
friend id operator*= (id &lhs, const size_t &rhs)
friend id operator/= (id &lhs, const size_t &rhs)
friend id operator%= (id &lhs, const size_t &rhs)
friend id operator<<= (id &lhs, const size_t &rhs)
friend id operator>>= (id &lhs, const size_t &rhs)
friend id operator&= (id &lhs, const size_t &rhs)
friend id operator|= (id &lhs, const size_t &rhs)
friend id operator^= (id &lhs, const size_t &rhs)

```

4.3.3.4 Item

A DPC++ `item` encapsulates a function object executing on an individual data point in a DPC++ range. When a kernel is executed, it is associated with an individual item in a range and acts upon it. This association is accomplished implicitly, by the runtime. Therefore, there are no user callable constructors; a DPC++ item is created when a kernel is instantiated.

The member functions of the item class pertain to determining the relationship between the item and the enclosing range:

- `get_id` – obtain the position of the work item in the iteration space
- `get_range` – obtain the range associated with the item
- `get_offset` – obtain the position of the item in the n-dimensional space

- `get_linear_id` – obtain the position of the item converting the n-dimensional space into one

```
template<int dimensions = 1, bool with_offset = true>  
class item
```

```
    tparam dimensions
```

```
    tparam with_offset
```

```
    id<dimensions> get_id() const
```

```
    size_t get_id(int dimension) const
```

```
    Parameters dimension –
```

```
    Returns
```

```
    size_t operator[] (int dimension) const
```

```
    range<dimensions> get_range() const
```

```
    size_t get_range (int dimension) const
```

```
    Parameters dimension –
```

```
    Returns
```

```
    id<dimensions> get_offset() const
```

```
    Returns
```

```
    operator item<dimensions, true>() const
```

```
    size_t get_linear_id() const
```

4.3.3.5 Nd_item

A DPC++ `nd_item` encapsulates a function object executing on an individual data point in a DPC++ `nd_range`. When a kernel is executed, it is associated with an individual item in a range and acts upon it. This association is accomplished implicitly, by the runtime. Therefore, there are no user callable constructors; a DPC++ `nd_item` is created when a kernel is instantiated.

The member functions of the `nd_item` class pertain to determining the relationship between the `nd_item` and the enclosing range:

- `get_global_id` – obtain the position of the work item in the iteration space
- `get_global_linear_id` – obtain the position of the work item in a linear representation of the global iteration space
- `get_local_id` – obtain the position of the item in the current work-group
- `get_local_linear_id` – obtain the position of the item in a linear representation of the current work-group
- `get_group` – obtain the position of the item in the overall `nd_range`
- `get_group_range` – obtain the number of work-groups in the iteration space
- `get_global_range` – obtain the range representing the dimensions of the global iteration space
- `get_local_range` – obtain the range representing the dimension of the current work-group
- `get_offset` – obtain an id that represents the offset between a work-item representation between local and global iteration space
- `get_nd_range` – obtain the `nd_range` from the `nd_item`

The class also includes a member function, `async_work_group_copy`, which can copy a range of items asynchronously.

```
template<int dimensions = 1>
class nd_item

    id<dimensions> get_global_id() const
    size_t get_global_id(int dimension) const
    size_t get_global_linear_id() const
    id<dimensions> get_local_id() const
    size_t get_local_id(int dimension) const
    size_t get_local_linear_id() const

    group<dimensions> get_group() const
    size_t get_group(int dimension) const
    size_t get_group_linear_id() const

    range<dimensions> get_group_range() const
    size_t get_group_range(int dimension) const

    range<dimensions> get_global_range() const
    size_t get_global_range(int dimension) const

    range<dimensions> get_local_range() const
    size_t get_local_range(int dimension) const

    id<dimensions> get_offset() const

    nd_range<dimensions> get_nd_range() const

    void barrier(access::fence_space accessSpace = access::fence_space::global_and_local) const

    template<access::mode accessMode = access::mode::read_write>
    void mem_fence(access::fence_space accessSpace = access::fence_space::global_and_local) const

    template<typename dataT>
    device_event async_work_group_copy(local_ptr<dataT> dest, global_ptr<dataT> src, size_t numElements) const

    template<typename dataT>
    device_event async_work_group_copy(global_ptr<dataT> dest, local_ptr<dataT> src, size_t numElements) const

    template<typename dataT>
    device_event async_work_group_copy(local_ptr<dataT> dest, global_ptr<dataT> src, size_t numElements, size_t srcStride) const

    template<typename dataT>
    device_event async_work_group_copy(global_ptr<dataT> dest, local_ptr<dataT> src, size_t numElements, size_t destStride) const

    template<typename ...eventTN>
    void wait_for(eventTN... events) const
```

4.3.3.6 Group

The `group` class encapsulates work-group functionality. Constructors are not user-callable; objects are created as a by-product of a call to `parallel_for_work_group`.

Once a group object has been instantiated, query various properties of the object by calling several member functions including:

- `get_id` – obtains the index of the work-group
- `get_global_range` – obtain a range that represents the work-items across the index space
- `get_local_range` – obtain a range that represents the work-items in a work-group
- `get_group_range` – obtain a range representing the dimensions of the current work-group
- `get_linear_id` – obtain a linear version of the work-group id

Definitions of global range and local range are in the SYCL Specification glossary. In brief, a global range is the overall number of work-items in the index space. A local range is the number of work-items in a work-group.

```
template<int dimensions = 1>
```

```
class group
```

```
    id<dimensions> get_id() const
    size_t get_id(int dimension) const

    range<dimensions> get_global_range() const
    size_t get_global_range(int dimension) const

    range<dimensions> get_local_range() const
    size_t get_local_range(int dimension) const

    range<dimensions> get_group_range() const
    size_t get_group_range(int dimension) const

    size_t operator[] (int dimension) const

    size_t get_linear_id() const

    template<typename workItemFunctionT>
    void parallel_for_work_item(workItemFunctionT func) const

    template<typename workItemFunctionT>
    void parallel_for_work_item(range<dimensions> logicalRange, workItemFunctionT func)
                                const

    template<access::mode accessMode = access::mode::read_write>
    void mem_fence (access::fence_space accessSpace = access::fence_space::global_and_local) const

    template<typename dataT>
    device_event async_work_group_copy (local_ptr<dataT> dest, global_ptr<dataT> src, size_t numElements) const

    template<typename dataT>
    device_event async_work_group_copy (global_ptr<dataT> dest, local_ptr<dataT> src, size_t numElements) const

    template<typename dataT>
    device_event async_work_group_copy (local_ptr<dataT> dest, global_ptr<dataT> src, size_t numElements, size_t srcStride) const

    template<typename dataT>
    device_event async_work_group_copy (global_ptr<dataT> dest, local_ptr<dataT> src, size_t numElements, size_t destStride) const
```

```
template<typename ...eventTN>
void wait_for(eventTN... events) const
```

4.3.3.7 Device Event

The DPC++ `device_event` class encapsulates wait objects within kernels. The `device_event` objects are used to coordinate asynchronous operations in kernels. The constructor and its arguments are unspecified and implementation dependent. The `wait` member function causes execution to stop until the operation associated with the wait is complete.

```
class device_event
```

```
void wait ()
```

4.3.3.8 Command Group Handler

The command group handler class encapsulates the actions of the command group, namely the marshaling of data and launching of the kernels on the devices.

There are no user callable constructors; construction is accomplished by the oneAPI runtime. Consider the example code below:

```
d_queue.submit([&](sycl::handler &cgh) {
    auto c_res = c_device.get_access<sycl::access::mode::write>(cgh);
    cgh.parallel_for<class ex1>(a_size, [=](sycl::id<1> idx) {
        c_res[idx] = 0;
    });
});
```

In the example, the accessor, `c_res`, is obtained from the device and takes a command group handler as a parameter, in this case `cgh`. The kernel dispatch itself is a member function of the command group handler. In this case, a `parallel_for` is called. The kernel dispatch API has multiple calls including `parallel_for`, `parallel_for_work_group`, and `single_task`.

There is a `set_args` function employed for passing arguments to an OpenCL™ kernel for interoperability.

```
class handler
```

```
template<typename dataT, int dimensions, access::mode accessMode, access::target accessTarget>
void require(accessor<dataT, dimensions, accessMode, accessTarget, access::placeholder::true_t>
             acc)
```

Template Parameters

- `dataT` –
- `dimensions` –
- `accessMode` –
- `accessTarget` –

OpenCL interoperability interface

```
template<typename T>
void set_arg (int argIndex, T &&arg)
```

Template Parameters **T** –

Parameters

- **argIndex** –
- **arg** –

```
template<typename ...Ts>
void set_args (Ts&&... args)
```

Kernel dispatch API

Note: Note: In all Kernel dispatch functions, when using a functor with a globally visible name the template parameter: "typename kernelName" can be omitted and the kernelType can be used instead.

```
template<typename KernelName, typename KernelType>
void single_task (KernelType kernelFunc)
void single_task (kernel syclKernel)
```

Template Parameters

- **KernelName** –
- **KernelType** –

Parameters

- **kernelFunc** –
- **syclKernel** –

```
template<typename KernelName, typename KernelType, int dimensions>
void parallel_for (range<dimensions> numWorkItems, KernelType kernelFunc)

template<typename KernelName, typename KernelType, int dimensions>
void parallel_for (range<dimensions> numWorkItems, id<dimensions> workItemOffset, KernelType
    kernelFunc)

template<typename KernelName, typename KernelType, int dimensions>
void parallel_for (nd_range<dimensions> executionRange, KernelType kernelFunc)

template<int dimensions>
void parallel_for (range<dimensions> numWorkItems, kernel syclKernel)

template<int dimensions>
void parallel_for (range<dimensions> numWorkItems, id<dimensions> workItemOffset, kernel
    syclKernel)

template<int dimensions>
void parallel_for (nd_range<dimensions> ndRange, kernel syclKernel)
```

Template Parameters

- **KernelName** –
- **KernelType** –
- **dimensions** –

Parameters

- **numWorkItems** –
- **workItemOffset** –
- **executionRange** –
- **ndRange** –
- **kernelFunc** –
- **syclKernel** –

```
template<typename KernelName, typename WorkgroupFunctionType, int dimensions>
void parallel_for_work_group (range<dimensions> numWorkGroups, WorkgroupFunctionType
                             kernelFunc)
```

```
template<typename KernelName, typename WorkgroupFunctionType, int dimensions>
void parallel_for_work_group (range<dimensions> numWorkGroups, range<dimensions> work-
                             GroupSize, WorkgroupFunctionType kernelFunc)
```

Template Parameters

- **KernelName** –
- **WorkgroupFunctionType** –
- **dimensions** –

Parameters

- **numWorkGroups** –
- **kernelFunc** –

Explicit memory operation APIs

```
template<typename T_src, int dim_src, access::mode mode_src, access::target tgt_src, access::placeholder isPlaceh
void copy (accessor<T_src, dim_src, mode_src, tgt_src, isPlaceholder> src, shared_ptr_class<T_dest>
           dest)
```

```
template<typename T_src, typename T_dest, int dim_dest, access::mode mode_dest, access::target tgt_dest, access:
void copy (shared_ptr_class<T_src> src, accessor<T_dest, dim_dest, mode_dest, tgt_dest, isPlace
           holder> dest)
```

```
template<typename T_src, int dim_src, access::mode mode_src, access::target tgt_src, access::placeholder isPlaceh
void copy (accessor<T_src, dim_src, mode_src, tgt_src, isPlaceholder> src, T_dest *dest)
```

```
template<typename T_src, typename T_dest, int dim_dest, access::mode mode_dest, access::target tgt_dest, access:
void copy (const T_src *src, accessor<T_dest, dim_dest, mode_dest, tgt_dest, isPlaceholder> dest)
```

```
template<typename T_src, int dim_src, access::mode mode_src, access::target tgt_src, access::placeholder isPlaceh
void copy (accessor<T_src, dim_src, mode_src, tgt_src, isPlaceholder_src> src, accessor<T_dest,
           dim_dest, mode_dest, tgt_dest, isPlaceholder_dest> dest)
```

Template Parameters

- **T_src** –
- **dim_src** –
- **mode_src** –
- **tgt_src** –
- **isPlaceholder** –
- **isPlaceholder_src** –

- **T_dest** –
- **dim_dest** –
- **mode_dest** –
- **tgt_dest** –
- **isPlaceholder_dest** –

Parameters

- **src** –
- **dest** –

```
template<typename T, int dim, access::mode mode, access::target tgt, access::placeholder isPlaceholder>  
void update_host (accessor<T, dim, mode, tgt, isPlaceholder> acc)
```

Template Parameters

- **T** –
- **dim** –
- **mode** –
- **tgt** –
- **isPlaceholder** –

Parameters *acc* –

```
template<typename T, int dim, access::mode mode, access::target tgt, access::placeholder isPlaceholder>  
void fill (accessor<T, dim, mode, tgt, isPlaceholder> dest, const T &src)
```

Template Parameters

- **T** –
- **dim** –
- **mode** –
- **tgt** –
- **isPlaceholder** –

Parameters

- **dest** –
- **src** –

4.3.3.9 Kernel

The DPC++ `kernel` class encapsulates methods and data for executing code on the device when a command group is instantiated. In many cases, the runtime creates the kernel objects when a command queue is instantiated.

Typically, a kernel object is not explicitly constructed by the user; instead it is constructed when a kernel dispatch function, such as `parallel_for`, is called. The sole case where a kernel object is constructed is when constructing a kernel object from an OpenCL application's `cl_kernel`. To compile the kernel ahead of time for use by the command queue, use the `program` class.

Member functions of the class return related objects and attributes regarding the kernel object including:

- `get` – obtains a `cl_kernel` if associated

- `Is_host` – obtains if the kernel is for the host
- `get_context` – obtains the context to which the kernel is associated
- `get_program` – obtains the program the kernel is contained in
- `get_info` – obtain details on the kernel and return in `info::kernel_info` descriptor
- `get_work_group_info` – obtain details on the work group and return in `info::kernel_work_group` descriptor

The `get_info` member function obtains kernel information such as `function_name`, `num_args`, `context`, `program`, `reference_count`, and `attributes`.

enum class kernel

```

    enumerator function_name
    enumerator num_args
    enumerator context
    enumerator program
    enumerator reference_count
    enumerator attributes

```

enum class kernel_work_group

```

    enumerator global_work_size
    enumerator work_group_size
    enumerator compile_work_group_size
    enumerator preferred_work_group_size_multiple
    enumerator private_mem_size

```

4.3.3.10 Program

A DPC++ program class encapsulates a program, either a host program or an OpenCL program. The program object is employed when compilation or linkage of the program is desired.

Constructors for a program object require a context at a minimum.

Member functions of the program class include:

- `get` – obtain an OpenCL program object from the program
- `is_host` – determines if the program is targeted for the host
- `compile_with_kernel_type` – enables compilation of a kernel
- `compile_with_source` – compiles OpenCL kernel
- `build_with_kernel_type` – builds kernel function
- `build_with_source` – builds kernel function from source
- `link` – link the object files
- `has_kernel` – determines if the program has a valid kernel function
- `get_kernel` – obtains the kernel from the program
- `get_binaries` – obtain a vector of compiled binaries for each device in the program

- `get_context` – obtain the context the program was built with
- `get_devices` – obtain a vector of the compiled binary sizes for each device

enum class program

```
    enumerator context
    enumerator devices
    enumerator reference_count
```

4.3.4 Error Handling

4.3.4.1 Exception

The DPC++ `exception` class encapsulates objects to communicate error conditions from the DPC++ program. Errors during DPC++ program execution are either scheduling or device related.

Execution between host and device is asynchronous in nature, therefore any events or errors are asynchronous. To catch exceptions that occur on the device, employ an `async_handler`, which is provided during command queue construction. During execution of the kernel, if any exceptions occur, these are placed on the `async_handler` list for processing once the command group function returns and the host handles the exceptions through the `async_handler` list. The `exception_ptr_class` is used to store the exception and can contain exceptions representing different types of errors such as `device_error`, `compile_program_error`, `link_program_error`, `invalid_object_error`, `memory_allocation_error`, `platform_error`, `profiling_error`, and `feature_not_supported`.

```
using async_handler = function_class<void (cl::sycl::exception_list) >
```

class exception

```
    const char *what () const
    bool has_context () const
    context get_context () const
    cl_int get_cl_code () const
```

class exception_list

Used as a container for a list of asynchronous exceptions

```
using value_type = exception_ptr_class
using reference = value_type&
using const_reference = const value_type&
using size_type = std::size_t
using iterator = __unspecified__
using const_iterator = __unspecified__

size_type size () const
iterator begin () const
    first asynchronous exception

iterator end () const
    refer to past-the-end last asynchronous exception
```

class runtime_error : public *exception*


```

class kernel_error : public runtime_error
class accessor_error : public runtime_error
class nd_range_error : public runtime_error
class event_error : public runtime_error
class invalid_parameter_error : public runtime_error
class device_error : public exception
class compile_program_error : public device_error
class link_program_error : public device_error
class invalid_object_error : public device_error
class memory_allocation_error : public device_error
class platform_error : public device_error
class profiling_error : public device_error
class feature_not_supported : public device_error

```

4.3.5 Stream

The DPC++ `stream` class is employed for outputting values of SYCL built-in, vector, and other types to the console.

```

enum class stream_manipulator

    enumerator flush
    enumerator dec
    enumerator hex
    enumerator oct
    enumerator noshowbase
    enumerator showbase
    enumerator noshowpos
    enumerator showpos
    enumerator endl
    enumerator fixed
    enumerator scientific
    enumerator hexfloat
    enumerator defaultfloat

const stream_manipulator flush = stream_manipulator::flush
const stream_manipulator dec = stream_manipulator::dec
const stream_manipulator hex = stream_manipulator::hex
const stream_manipulator oct = stream_manipulator::oct
const stream_manipulator noshowbase = stream_manipulator::noshowbase
const stream_manipulator showbase = stream_manipulator::showbase
const stream_manipulator noshowpos = stream_manipulator::noshowpos

```

```
const stream_manipulator showpos = stream_manipulator::showpos
const stream_manipulator endl = stream_manipulator::endl
const stream_manipulator fixed = stream_manipulator::fixed
const stream_manipulator scientific = stream_manipulator::scientific
const stream_manipulator hexfloat = stream_manipulator::hexfloat
const stream_manipulator defaultfloat = stream_manipulator::defaultfloat
__precision_manipulator__ setprecision (int precision)
__width_manipulator__ setw (int width)
class stream

    stream (size_t totalBufferSize, size_t workItemBufferSize, handler &cgh)
    size_t get_size () const
    size_t get_work_item_buffer_size () const
    size_t get_max_statement_size () const
    get_max_statement_size () has the same functionality as
    get_work_item_buffer_size (), and is provided for backward compatibility.
    get_max_statement_size() is a deprecated query.
```

4.3.6 Vec and Swizzled Vec

The DPC++ Vec and Swizzled Vec class templates are designed to represent vectors between host and devices.

To instantiate a `vec` class template, provide the type and an integer representing the number of elements. The number of elements can be 1, 2, 3, 4, 8, or 16; any other integer results in a compile error. The type provided must be a basic scalar type, such as `int` or `float`.

Member functions once an object is created include:

- `get_count` – obtains the number of elements of the `vec`
- `get_size` – obtain the size of the `vec` (in bytes)
- `lo` – obtain the lower half of the `vec`
- `hi` – obtain the higher half of the `vec`
- `odd` – obtain the odd index elements of the `vec`
- `even` – obtain the even index elements of the `vec`
- `load` – copy the pointed to values into a `vec`
- `store` – copy the `vec` into the pointed to location

The `__swizzled_vec__` class is employed to reposition elements of a `vec` object. A good motivation for employing is to obtain every odd or even element of a vector. In this case, employ the `odd` or `even` member function of the class. There are member functions associated with the `__swizzled_vec__` class for converting a `vec` into a new `vec`, such as one in `RGBA` format.

Various operators on the `vec` class include: `+=`, `-=`, `*=`, `/=`, `%=`, `++`, `--`, `&`, `l`, `^`, `+`, `-`, `*`, `/`, `%`, `&&`, `||`, `<<`, `>>`, `<=<=`, `>>=`, `==`, `!=`, `<`, `>`, `<=`, `>=`.

```
enum class rounding_mode
```

```

    enumerator automatic
    enumerator rte
    enumerator rtz
    enumerator rtp
    enumerator rtn

struct elem

    static constexpr int x = 0
    static constexpr int y = 1
    static constexpr int z = 2
    static constexpr int w = 3
    static constexpr int r = 0
    static constexpr int g = 1
    static constexpr int b = 2
    static constexpr int a = 3
    static constexpr int s0 = 0
    static constexpr int s1 = 1
    static constexpr int s2 = 2
    static constexpr int s3 = 3
    static constexpr int s4 = 4
    static constexpr int s5 = 5
    static constexpr int s6 = 6
    static constexpr int s7 = 7
    static constexpr int s8 = 8
    static constexpr int s9 = 9
    static constexpr int sA = 10
    static constexpr int sB = 11
    static constexpr int sC = 12
    static constexpr int sD = 13
    static constexpr int sE = 14
    static constexpr int sF = 15

template<typename dataT, int numElements>
class vec

    using element_type = dataT

    vec ()
    explicit vec (const dataT &arg)
    template<typename ...argTN>

```

```

vec (const argTN&... args)
vec (const vec<dataT, numElements> &rhs)
template<typename convertT, rounding_mode roundingMode = rounding_mode::automatic>
vec<convertT, numElements> convert () const

operator dataT () const
    Available only when: numElements == 1

size_t get_count () const

size_t get_size () const

template<typename asT>
asT as () const

template<int... swizzleIndexes>
__swizzled_vec__ swizzle () const

__swizzled_vec__ x_ACCESS () const
__swizzled_vec__ y_ACCESS () const
__swizzled_vec__ z_ACCESS () const
__swizzled_vec__ w_ACCESS () const
    Available only when numElements <= 4

__swizzled_vec__ r () const
__swizzled_vec__ g () const
__swizzled_vec__ b () const
__swizzled_vec__ a () const
    Available only numElements == 4

__swizzled_vec__ s0 () const
__swizzled_vec__ s1 () const
__swizzled_vec__ s2 () const
__swizzled_vec__ s3 () const
__swizzled_vec__ s4 () const
__swizzled_vec__ s5 () const
__swizzled_vec__ s6 () const
__swizzled_vec__ s7 () const
__swizzled_vec__ s8 () const
__swizzled_vec__ s9 () const
__swizzled_vec__ sA () const
__swizzled_vec__ sB () const
__swizzled_vec__ sC () const
__swizzled_vec__ sD () const
__swizzled_vec__ sE () const
__swizzled_vec__ sF () const

__swizzled_vec__ lo () const
__swizzled_vec__ hi () const
__swizzled_vec__ odd () const
__swizzled_vec__ even () const

template<access::address_space addressSpace>

```

```

void load (size_t offset, multi_ptr<const dataT, addressSpace> ptr)
template<access::address_space addressSpace>
void store (size_t offset, multi_ptr<dataT, addressSpace> ptr) const

    load and store member functions

friend vec operatorOP (const vec &lhs, const vec &rhs)
friend vec operatorOP (const vec &lhs, const dataT &rhs)
    OP is: +, -, *, /, %

friend vec &operatorOP (vec &lhs, const vec &rhs)
friend vec &operatorOP (vec &lhs, const dataT &rhs)
    OP is: +=, -=, *=, /=, %=

friend vec &operatorOP (vec &lhs)
friend vec operatorOP (vec &lhs, int)
    OP is: ++, --

friend vec operator& (const vec &lhs, const vec &rhs)
friend vec operator& (const vec &lhs, const dataT &rhs)
friend vec operator| (const vec &lhs, const vec &rhs)
friend vec operator| (const vec &lhs, const dataT &rhs)
friend vec operator^ (const vec &lhs, const vec &rhs)
friend vec operator^ (const vec &lhs, const dataT &rhs)

friend vec &operator&= (vec &lhs, const vec &rhs)
friend vec &operator&= (vec &lhs, const dataT &rhs)
friend vec &operator|= (vec &lhs, const vec &rhs)
friend vec &operator|= (vec &lhs, const dataT &rhs)
friend vec &operator^= (vec &lhs, const vec &rhs)
friend vec &operator^= (vec &lhs, const dataT &rhs)

friend vec<RET, numElements> operator&& (const vec &lhs, const vec &rhs)
friend vec<RET, numElements> operator&& (const vec &lhs, const dataT &rhs)
friend vec<RET, numElements> operator|| (const vec &lhs, const vec &rhs)
friend vec<RET, numElements> operator|| (const vec &lhs, const dataT &rhs)

friend vec operator<< (const vec &lhs, const vec &rhs)
friend vec operator<< (const vec &lhs, const dataT &rhs)
friend vec operator>> (const vec &lhs, const vec &rhs)
friend vec operator>> (const vec &lhs, const dataT &rhs)

friend vec &operator<<= (vec &lhs, const vec &rhs)
friend vec &operator<<= (vec &lhs, const dataT &rhs)
friend vec &operator>>= (vec &lhs, const vec &rhs)
friend vec &operator>>= (vec &lhs, const dataT &rhs)

friend vec<RET, numElements> operator== (const vec &lhs, const vec &rhs)
friend vec<RET, numElements> operator== (const vec &lhs, const dataT &rhs)
friend vec<RET, numElements> operator!= (const vec &lhs, const vec &rhs)
friend vec<RET, numElements> operator!= (const vec &lhs, const dataT &rhs)
friend vec<RET, numElements> operator< (const vec &lhs, const vec &rhs)

```

```
friend vec<RET, numElements> operator< (const vec &lhs, const dataT &rhs)
friend vec<RET, numElements> operator> (const vec &lhs, const vec &rhs)
friend vec<RET, numElements> operator> (const vec &lhs, const dataT &rhs)
friend vec<RET, numElements> operator<= (const vec &lhs, const vec &rhs)
friend vec<RET, numElements> operator<= (const vec &lhs, const dataT &rhs)
friend vec<RET, numElements> operator>= (const vec &lhs, const vec &rhs)
friend vec<RET, numElements> operator>= (const vec &lhs, const dataT &rhs)
vec<dataT, numElements> &operator= (const vec<dataT, numElements> &rhs)
vec<dataT, numElements> &operator= (const dataT &rhs)

friend vec operator~ (const vec &v)
friend vec<RET, numElements> operator! (const vec &v)

friend vec operatorOP (const dataT &lhs, const vec &rhs)
    OP is: +, -, *, /, %

friend vec operator& (const dataT &lhs, const vec &rhs)
friend vec operator| (const dataT &lhs, const vec &rhs)
friend vec operator^ (const dataT &lhs, const vec &rhs)

friend vec<RET, numElements> operator&& (const dataT &lhs, const vec &rhs)
friend vec<RET, numElements> operator|| (const dataT &lhs, const vec &rhs)

friend vec operator<< (const dataT &lhs, const vec &rhs)
friend vec operator>> (const dataT &lhs, const vec &rhs)

friend vec<RET, numElements> operator== (const dataT &lhs, const vec &rhs)
friend vec<RET, numElements> operator!= (const dataT &lhs, const vec &rhs)
friend vec<RET, numElements> operator< (const dataT &lhs, const vec &rhs)
friend vec<RET, numElements> operator> (const dataT &lhs, const vec &rhs)
friend vec<RET, numElements> operator<= (const dataT &lhs, const vec &rhs)
friend vec<RET, numElements> operator>= (const dataT &lhs, const vec &rhs)
```

4.3.7 Built-in Types & Functions

The DPC++ built-in functions provide low level capabilities that can execute on the host and device with some level of compatibility. Section 4.13 of the SYCL Specification details all the various built-in types and functions available.

One task taken care of by the implementation is the mapping of C++ fundamental types such as int, short, long such that the types agree between the host and the device.

The built-in scalar data types are summarized in the SYCL Specification. In general, the built-in types cover floating point, double, half precision, char, signed char, unsigned char, short, unsigned short, int, unsigned int, long, unsigned long, long long, unsigned long long, and signed integer. Lastly, the built-in types can be post fixed with 2, 3, 4, 8, or 16, which indicated a vector of the post fixed type. For example, float3 indicates a type consisting of three floating point elements addressed as one object. A float3 is common in image processing algorithms for representing RGB data.

The built-in functions are either defined as part of lower level classes or part of the basic math functions. These built-in functions enable vendors to provide differentiated built-in functions specific to the architecture while also enabling basic functionality for generic implementations.

The categories of built-in functions are summarized as:

- Work-item functions – pertaining to `nd_item` and `group` classes
- Basic Math functions – low level math and comparison functions
- General math functions – Transcendental, trigonometric, and geometric functions
- Vector load and store – reading and writing `vec` class
- Synchronization – `nd_item` related barriers
- Output – `stream` class for output

4.3.8 Property Interface

The DPC++ property interface is employed with the `buffer`, `image`, and `queue` classes to provide extra information to those classes without affecting the type. These classes provide an additional `has_property` and `get_property` member function to test for and obtain a particular property.

```
template<typename propertyT>
struct is_property

template<typename propertyT, typename syclObjectT>
struct is_property_of

class T

    template<typename propertyT>
    bool has_property() const

    template<typename propertyT>
    propertyT get_property() const

class property_list

    template<typename ...propertyTN>
    property_list (propertyTN... props)
```

4.3.9 Version

The include file, `version.h`, includes a definition of `__SYCL_COMPILER_VERSION` based upon the date of the compiler. It can be used to control compilation based upon the specific version of the compiler.

GLOSSARY

accelerator Specialized component containing compute resources that can quickly execute a subset of operations. Examples include CPU, FPGA, GPU. See also: *device*

accessor Communicates the desired location (host, device) and mode (read, write) of access.

application scope Code that executes on the host.

buffers Memory object that communicates the type and number of items of that type to be communicated to the device for computation.

command group scope Code that acts as the interface between the host and device.

command queue Issues command groups concurrently.

compute unit A grouping of processing elements into a ‘core’ that contains shared elements for use between the processing elements and with faster access than memory residing on other compute units on the device.

device An accelerator or specialized component containing compute resources that can quickly execute a subset of operations. A CPU can be employed as a device, but when it is, it is being employed as an accelerator. Examples include CPU, FPGA, GPU. See also: *accelerator*

device code Code that executes on the device rather than the host. Device code is specified via lambda expression, functor, or kernel class.

fat binary Application binary that contains device code for multiple devices. The binary includes both the generic code (SPIR-V representation) and target specific executable code.

fat library Archive or library of object code that contains object code for multiple devices. The fat library includes both the generic object (SPIR-V representation) and target specific object code.

fat object File that contains object code for multiple devices. The fat object includes both the generic object (SPIR-V representation) and target specific object code.

host A CPU-based system (computer) that executes the primary portion of a program, specifically the application scope and command group scope.

host device A SYCL device that is always present and usually executes on the host CPU.

host code Code that is compiled by the host compiler and executes on the host rather than the device.

images Formatted opaque memory object that is accessed via built-in function. Typically pertains to pictures comprised of pixels stored in format like RGB.

kernel scope Code that executes on the device.

nd-range Short for N-Dimensional Range, a group of kernel instances, or work item, across one, two, or three dimensions.

processing element Individual engine for computation that makes up a compute unit.

single source Code in the same file that can execute on a host and accelerator(s).

SPIR-V Binary intermediate language for representing graphical-shader stages and compute kernels.

sub-group Sub-groups are an Intel extension.

work-group Collection of work-items that execute on a compute unit.

work-item Basic unit of computation in the oneAPI programming model. It is associated with a kernel which executes on the processing element.

NOTICES AND DISCLAIMERS

No license (express or implied, by estoppel or otherwise) to any intellectual property rights is granted by this document.

Intel technologies' features and benefits depend on system configuration and may require enabled hardware, software or service activation. Performance varies depending on system configuration. No product or component can be absolutely secure. Check with your system manufacturer or retailer or learn more at [intel.com].

Intel disclaims all express and implied warranties, including without limitation, the implied warranties of merchantability, fitness for a particular purpose, and non-infringement, as well as any warranty arising from course of performance, course of dealing, or usage in trade.

This document contains information on products, services and/or processes in development. All information provided here is subject to change without notice. Contact your Intel representative to obtain the latest forecast, schedule, specifications and roadmaps.

The products and services described may contain defects or errors which may cause deviations from published specifications. Current characterized errata are available on request.

Intel, the Intel logo, Intel Atom, Intel Core, Intel Xeon Phi, VTune and Xeon are trademarks of Intel Corporation in the U.S. and/or other countries.

*Other names and brands may be claimed as the property of others.

Microsoft, Windows, and the Windows logo are trademarks, or registered trademarks of Microsoft Corporation in the United States and/or other countries.

OpenCL and the OpenCL logo are trademarks of Apple Inc. used by permission by Khronos.

Java is a registered trademark of Oracle and/or its affiliates.

© Intel Corporation.

This software and the related documents are Intel copyrighted materials, and your use of them is governed by the express license under which they were provided to you (**License**). Unless the License provides otherwise, you may not use, modify, copy, publish, distribute, disclose or transmit this software or the related documents without Intel's prior written permission.

This software and the related documents are provided as is, with no express or implied warranties, other than those that are expressly stated in the License.

| |
|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Intel's compilers may or may not optimize to the same degree for non-Intel microprocessors for optimizations that are not unique to Intel microprocessors. These optimizations include SSE2, SSE3, and SSSE3 instruction sets and other optimizations. Intel does not guarantee the availability, functionality, or effectiveness of any optimization on microprocessors not manufactured by Intel. Microprocessor-dependent optimizations in this product are intended for use with Intel microprocessors. Certain optimizations not specific to Intel microarchitecture are reserved for Intel microprocessors. Please refer to the applicable product User and Reference Guides for more information regarding the specific instruction sets covered by this notice. Notice revision #20110804 |
|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|

Unless stated otherwise, the code examples in this document are provided to you under an MIT license, the terms of which are as follows:

Copyright 2019 Intel Corporation

Permission is hereby granted, free of charge, to any person obtaining a copy of this software and associated documentation files (the “Software”), to deal in the Software without restriction, including without limitation the rights to use, copy, modify, merge, publish, distribute, sublicense, and/or sell copies of the Software, and to permit persons to whom the Software is furnished to do so, subject to the following conditions:

The above copyright notice and this permission notice shall be included in all copies or substantial portions of the Software.

THE SOFTWARE IS PROVIDED “AS IS”, WITHOUT WARRANTY OF ANY KIND, EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO THE WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE AND NONINFRINGEMENT. IN NO EVENT SHALL THE AUTHORS OR COPYRIGHT HOLDERS BE LIABLE FOR ANY CLAIM, DAMAGES OR OTHER LIABILITY, WHETHER IN AN ACTION OF CONTRACT, TORT OR OTHERWISE, ARISING FROM, OUT OF OR IN CONNECTION WITH THE SOFTWARE OR THE USE OR OTHER DEALINGS IN THE SOFTWARE.

A

accelerator, [69](#)
 access::address_space (C++ enum), [44](#)
 access::address_space::constant_space (C++ enumerator), [44](#)
 access::address_space::global_space (C++ enumerator), [44](#)
 access::address_space::local_space (C++ enumerator), [44](#)
 access::address_space::private_space (C++ enumerator), [44](#)
 accessor, [69](#)
 accessor_error (C++ class), [61](#)
 application scope, [69](#)
 async_handler (C++ type), [60](#)
 atomic (C++ class), [39](#)
 atomic::atomic (C++ function), [39](#)
 atomic::compare_exchange_strong (C++ function), [39](#)
 atomic::exchange (C++ function), [39](#)
 atomic::fetch_add (C++ function), [40](#)
 atomic::fetch_and (C++ function), [40](#)
 atomic::fetch_max (C++ function), [40](#)
 atomic::fetch_min (C++ function), [40](#)
 atomic::fetch_or (C++ function), [40](#)
 atomic::fetch_sub (C++ function), [40](#)
 atomic::fetch_xor (C++ function), [40](#)
 atomic::load (C++ function), [39](#)
 atomic::store (C++ function), [39](#)

B

buffer (C++ class), [40](#)
 buffer::allocator_type (C++ type), [41](#)
 buffer::buffer (C++ function), [41](#)
 buffer::buffer<T, 1> (C++ function), [41](#)
 buffer::const_reference (C++ type), [41](#)
 buffer::get_access (C++ function), [42](#)
 buffer::get_allocator (C++ function), [42](#)
 buffer::get_count (C++ function), [42](#)
 buffer::get_range (C++ function), [42](#)
 buffer::get_size (C++ function), [42](#)
 buffer::is_sub_buffer (C++ function), [42](#)

buffer::reference (C++ type), [41](#)
 buffer::reinterpret (C++ function), [43](#)
 buffer::set_final_data (C++ function), [42](#)
 buffer::set_write_back (C++ function), [42](#)
 buffer::value_type (C++ type), [41](#)
 buffers, [69](#)

C

command group scope, [69](#)
 command queue, [69](#)
 compile_program_error (C++ class), [61](#)
 compute unit, [69](#)
 constant_ptr (C++ type), [46](#)

D

dec (C++ member), [61](#)
 defaultfloat (C++ member), [62](#)
 device, [69](#)
 device code, [69](#)
 device_error (C++ class), [61](#)
 device_event (C++ class), [55](#)
 device_event::wait (C++ function), [55](#)

E

elem (C++ struct), [63](#)
 elem::a (C++ member), [63](#)
 elem::b (C++ member), [63](#)
 elem::g (C++ member), [63](#)
 elem::r (C++ member), [63](#)
 elem::s0 (C++ member), [63](#)
 elem::s1 (C++ member), [63](#)
 elem::s2 (C++ member), [63](#)
 elem::s3 (C++ member), [63](#)
 elem::s4 (C++ member), [63](#)
 elem::s5 (C++ member), [63](#)
 elem::s6 (C++ member), [63](#)
 elem::s7 (C++ member), [63](#)
 elem::s8 (C++ member), [63](#)
 elem::s9 (C++ member), [63](#)
 elem::sA (C++ member), [63](#)
 elem::sB (C++ member), [63](#)
 elem::sC (C++ member), [63](#)

elem::sD (C++ member), 63
 elem::sE (C++ member), 63
 elem::sF (C++ member), 63
 elem::w (C++ member), 63
 elem::x (C++ member), 63
 elem::y (C++ member), 63
 elem::z (C++ member), 63
 endl (C++ member), 62
 event_error (C++ class), 61
 exception (C++ class), 60
 exception::get_cl_code (C++ function), 60
 exception::get_context (C++ function), 60
 exception::has_context (C++ function), 60
 exception::what (C++ function), 60
 exception_list (C++ class), 60
 exception_list::begin (C++ function), 60
 exception_list::const_iterator (C++ type), 60
 exception_list::const_reference (C++ type), 60
 exception_list::end (C++ function), 60
 exception_list::iterator (C++ type), 60
 exception_list::reference (C++ type), 60
 exception_list::size (C++ function), 60
 exception_list::size_type (C++ type), 60
 exception_list::value_type (C++ type), 60
 exception_ptr_class (C++ type), 22

F

fat binary, 69
 fat library, 69
 fat object, 69
 feature_not_supported (C++ class), 61
 fixed (C++ member), 62
 flush (C++ member), 61
 function_class (C++ type), 22

G

global_ptr (C++ type), 46
 group (C++ class), 54
 group::async_work_group_copy (C++ function), 54
 group::get_global_range (C++ function), 54
 group::get_group_range (C++ function), 54
 group::get_id (C++ function), 54
 group::get_linear_id (C++ function), 54
 group::get_local_range (C++ function), 54
 group::mem_fence (C++ function), 54
 group::operator[] (C++ function), 54
 group::parallel_for_work_item (C++ function), 54
 group::wait_for (C++ function), 54

H

handler (C++ class), 55
 handler::copy (C++ function), 57
 handler::fill (C++ function), 58
 handler::parallel_for (C++ function), 56
 handler::parallel_for_work_group (C++ function), 57
 handler::require (C++ function), 55
 handler::set_arg (C++ function), 56
 handler::set_args (C++ function), 56
 handler::single_task (C++ function), 56
 handler::update_host (C++ function), 58
 hash_class (C++ type), 22
 hex (C++ member), 61
 hexfloat (C++ member), 62
 host, 69
 host code, 69
 host device, 69

I

id (C++ class), 49
 id::get (C++ function), 50
 id::id (C++ function), 49
 id::operator* (C++ function), 50
 id::operator*= (C++ function), 51
 id::operator+ (C++ function), 50
 id::operator+= (C++ function), 51
 id::operator/ (C++ function), 50
 id::operator/= (C++ function), 51
 id::operator% (C++ function), 50
 id::operator%= (C++ function), 51
 id::operator& (C++ function), 50
 id::operator&= (C++ function), 51
 id::operator&& (C++ function), 50
 id::operator- (C++ function), 50
 id::operator^ (C++ function), 50
 id::operator^= (C++ function), 51
 id::operator| (C++ function), 50
 id::operator|= (C++ function), 51
 id::operator|| (C++ function), 50
 id::operator> (C++ function), 50
 id::operator>= (C++ function), 50
 id::operator>> (C++ function), 50
 id::operator>>= (C++ function), 51
 id::operator< (C++ function), 50
 id::operator<= (C++ function), 50
 id::operator<< (C++ function), 50
 id::operator<<= (C++ function), 51
 id::operator[] (C++ function), 50
 images, 69
 info::kernel (C++ enum), 59
 info::kernel::attributes (C++ enumerator), 59
 info::kernel::context (C++ enumerator), 59

[info::kernel::function_name \(C++ enumerator\), 59](#)
[info::kernel::num_args \(C++ enumerator\), 59](#)
[info::kernel::program \(C++ enumerator\), 59](#)
[info::kernel::reference_count \(C++ enumerator\), 59](#)
[info::kernel_work_group \(C++ enum\), 59](#)
[info::kernel_work_group::compile_work_group_size \(C++ enumerator\), 59](#)
[info::kernel_work_group::global_work_size \(C++ enumerator\), 59](#)
[info::kernel_work_group::preferred_work_group_size \(C++ enumerator\), 59](#)
[info::kernel_work_group::private_mem_size \(C++ enumerator\), 59](#)
[info::kernel_work_group::work_group_size \(C++ enumerator\), 59](#)
[info::program \(C++ enum\), 60](#)
[info::program::context \(C++ enumerator\), 60](#)
[info::program::devices \(C++ enumerator\), 60](#)
[info::program::reference_count \(C++ enumerator\), 60](#)
[invalid_object_error \(C++ class\), 61](#)
[invalid_parameter_error \(C++ class\), 61](#)
[is_property \(C++ struct\), 67](#)
[is_property_of \(C++ struct\), 67](#)
[item \(C++ class\), 52](#)
[item::get_id \(C++ function\), 52](#)
[item::get_linear_id \(C++ function\), 52](#)
[item::get_offset \(C++ function\), 52](#)
[item::get_range \(C++ function\), 52](#)
[item::operator item<dimensions, true> \(C++ function\), 52](#)
[item::operator\[\] \(C++ function\), 52](#)

K

[kernel scope, 69](#)
[kernel_error \(C++ class\), 60](#)

L

[link_program_error \(C++ class\), 61](#)
[local_ptr \(C++ type\), 46](#)

M

[memory_allocation_error \(C++ class\), 61](#)
[memory_order \(C++ enum\), 39](#)
[memory_order::relaxed \(C++ enumerator\), 39](#)
[multi_ptr \(C++ class\), 44](#)
[multi_ptr::~~multi_ptr \(C++ function\), 45](#)
[multi_ptr::address_space \(C++ member\), 44](#)
[multi_ptr::const_pointer_t \(C++ type\), 44](#)
[multi_ptr::const_reference_t \(C++ type\), 44](#)
[multi_ptr::difference_type \(C++ type\), 44](#)
[multi_ptr::element_type \(C++ type\), 44](#)
[multi_ptr::get \(C++ function\), 45](#)
[multi_ptr::multi_ptr \(C++ function\), 44](#)
[multi_ptr::operator ElementType* \(C++ function\), 45](#)
[multi_ptr::operator multi_ptr<const ElementType, Space> \(C++ function\), 45](#)
[multi_ptr::operator multi_ptr<const void, Space> \(C++ function\), 45](#)
[multi_ptr::operator multi_ptr<void, Space> \(C++ function\), 45](#)
[multi_ptr::operator!= \(C++ function\), 46](#)
[multi_ptr::operator* \(C++ function\), 45](#)
[multi_ptr::operator+ \(C++ function\), 45](#)
[multi_ptr::operator++ \(C++ function\), 45](#)
[multi_ptr::operator+= \(C++ function\), 45](#)
[multi_ptr::operator= \(C++ function\), 45](#)
[multi_ptr::operator== \(C++ function\), 46](#)
[multi_ptr::operator- \(C++ function\), 45](#)
[multi_ptr::operator-= \(C++ function\), 45](#)
[multi_ptr::operator-- \(C++ function\), 45](#)
[multi_ptr::operator-> \(C++ function\), 45](#)
[multi_ptr::operator> \(C++ function\), 46](#)
[multi_ptr::operator>= \(C++ function\), 46](#)
[multi_ptr::operator< \(C++ function\), 46](#)
[multi_ptr::operator<= \(C++ function\), 46](#)
[multi_ptr::pointer_t \(C++ type\), 44](#)
[multi_ptr::prefetch \(C++ function\), 45](#)
[multi_ptr::reference_t \(C++ type\), 44](#)
[mutex_class \(C++ type\), 22](#)

N

[nd_item \(C++ class\), 53](#)
[nd_item::async_work_group_copy \(C++ function\), 53](#)
[nd_item::barrier \(C++ function\), 53](#)
[nd_item::get_global_id \(C++ function\), 53](#)
[nd_item::get_global_linear_id \(C++ function\), 53](#)
[nd_item::get_global_range \(C++ function\), 53](#)
[nd_item::get_group \(C++ function\), 53](#)
[nd_item::get_group_linear_id \(C++ function\), 53](#)
[nd_item::get_group_range \(C++ function\), 53](#)
[nd_item::get_local_id \(C++ function\), 53](#)
[nd_item::get_local_linear_id \(C++ function\), 53](#)
[nd_item::get_local_range \(C++ function\), 53](#)
[nd_item::get_nd_range \(C++ function\), 53](#)
[nd_item::get_offset \(C++ function\), 53](#)
[nd_item::mem_fence \(C++ function\), 53](#)
[nd_item::wait_for \(C++ function\), 53](#)
[nd_range \(C++ class\), 49](#)

nd_range::get_global_range (C++ function), 49
 nd_range::get_group_range (C++ function), 49
 nd_range::get_local_range (C++ function), 49
 nd_range::get_offset (C++ function), 49
 nd_range::nd_range (C++ function), 49
 nd_range_error (C++ class), 61
 nd-range, 69
 noshowbase (C++ member), 61
 noshowpos (C++ member), 61

O

oct (C++ member), 61

P

platform_error (C++ class), 61
 private_ptr (C++ type), 46
 processing element, 69
 profiling_error (C++ class), 61
 property::buffer::context_bound (C++ class), 43
 property::buffer::context_bound::context_bound (C++ function), 43
 property::buffer::context_bound::get_context_bound (C++ function), 43
 property::buffer::use_host_ptr (C++ class), 43
 property::buffer::use_host_ptr::use_host_ptr (C++ function), 43
 property::buffer::use_mutex (C++ class), 43
 property::buffer::use_mutex::get_mutex_ptr (C++ function), 43
 property::buffer::use_mutex::use_mutex (C++ function), 43
 property_list (C++ class), 67
 property_list::property_list (C++ function), 67

R

range (C++ class), 46
 range::get (C++ function), 47
 range::operator* (C++ function), 47, 48
 range::operator*= (C++ function), 48
 range::operator+ (C++ function), 47, 48
 range::operator+= (C++ function), 48
 range::operator/ (C++ function), 47, 48
 range::operator/= (C++ function), 48
 range::operator% (C++ function), 47, 48
 range::operator%= (C++ function), 48
 range::operator& (C++ function), 47, 48
 range::operator&= (C++ function), 48
 range::operator&& (C++ function), 47, 48
 range::operator- (C++ function), 47, 48
 range::operator-= (C++ function), 48

range::operator^ (C++ function), 47, 48
 range::operator^= (C++ function), 48
 range::operator| (C++ function), 47, 48
 range::operator|= (C++ function), 48
 range::operator|| (C++ function), 47, 48
 range::operator> (C++ function), 47, 48
 range::operator>= (C++ function), 47, 48
 range::operator>> (C++ function), 47, 48
 range::operator>>= (C++ function), 48
 range::operator< (C++ function), 47, 48
 range::operator<= (C++ function), 47, 48
 range::operator<< (C++ function), 47, 48
 range::operator<<= (C++ function), 48
 range::operator[] (C++ function), 47
 range::range (C++ function), 47
 range::size (C++ function), 47
 rounding_mode (C++ enum), 62
 rounding_mode::automatic (C++ enumerator), 62
 rounding_mode::rte (C++ enumerator), 62
 rounding_mode::rtn (C++ enumerator), 62
 rounding_mode::rtp (C++ enumerator), 62
 rounding_mode::rtz (C++ enumerator), 62
 runtime_error (C++ class), 60

S

scientific (C++ member), 62
 setprecision (C++ function), 62
 setw (C++ function), 62
 shared_ptr_class (C++ type), 22
 showbase (C++ member), 61
 showpos (C++ member), 61
 single source, 70
 SPIR-V, 70
 stream (C++ class), 62
 stream::get_max_statement_size (C++ function), 62
 stream::get_size (C++ function), 62
 stream::get_work_item_buffer_size (C++ function), 62
 stream::stream (C++ function), 62
 stream_manipulator (C++ enum), 61
 stream_manipulator::dec (C++ enumerator), 61
 stream_manipulator::defaultfloat (C++ enumerator), 61
 stream_manipulator::endl (C++ enumerator), 61
 stream_manipulator::fixed (C++ enumerator), 61
 stream_manipulator::flush (C++ enumerator), 61
 stream_manipulator::hex (C++ enumerator), 61

stream_manipulator::hexfloat (C++ *enumerator*), 61
 stream_manipulator::noshowbase (C++ *enumerator*), 61
 stream_manipulator::noshowpos (C++ *enumerator*), 61
 stream_manipulator::oct (C++ *enumerator*), 61
 stream_manipulator::scientific (C++ *enumerator*), 61
 stream_manipulator::showbase (C++ *enumerator*), 61
 stream_manipulator::showpos (C++ *enumerator*), 61
 string_class (C++ *type*), 22
 sub-group, 70

T

T (C++ *class*), 67
 T::get_property (C++ *function*), 67
 T::has_property (C++ *function*), 67

U

unique_ptr_class (C++ *type*), 22

V

vec (C++ *class*), 63
 vec::a (C++ *function*), 64
 vec::as (C++ *function*), 64
 vec::b (C++ *function*), 64
 vec::convert (C++ *function*), 63
 vec::element_type (C++ *type*), 63
 vec::even (C++ *function*), 64
 vec::g (C++ *function*), 64
 vec::get_count (C++ *function*), 64
 vec::get_size (C++ *function*), 64
 vec::hi (C++ *function*), 64
 vec::lo (C++ *function*), 64
 vec::load (C++ *function*), 64
 vec::odd (C++ *function*), 64
 vec::operator dataT (C++ *function*), 64
 vec::operator! (C++ *function*), 66
 vec::operator!= (C++ *function*), 65, 66
 vec::operator= (C++ *function*), 66
 vec::operator== (C++ *function*), 65, 66
 vec::operator& (C++ *function*), 65, 66
 vec::operator&= (C++ *function*), 65
 vec::operator&& (C++ *function*), 65, 66
 vec::operator^ (C++ *function*), 65, 66
 vec::operator^= (C++ *function*), 65
 vec::operator~ (C++ *function*), 66
 vec::operator| (C++ *function*), 65, 66
 vec::operator|= (C++ *function*), 65
 vec::operator|| (C++ *function*), 65, 66
 vec::operator> (C++ *function*), 65, 66
 vec::operator>= (C++ *function*), 65, 66
 vec::operator>> (C++ *function*), 65, 66
 vec::operator>>= (C++ *function*), 65
 vec::operator< (C++ *function*), 65, 66
 vec::operator<= (C++ *function*), 65, 66
 vec::operator<< (C++ *function*), 65, 66
 vec::operator<<= (C++ *function*), 65
 vec::operatorOP (C++ *function*), 65, 66
 vec::r (C++ *function*), 64
 vec::s0 (C++ *function*), 64
 vec::s1 (C++ *function*), 64
 vec::s2 (C++ *function*), 64
 vec::s3 (C++ *function*), 64
 vec::s4 (C++ *function*), 64
 vec::s5 (C++ *function*), 64
 vec::s6 (C++ *function*), 64
 vec::s7 (C++ *function*), 64
 vec::s8 (C++ *function*), 64
 vec::s9 (C++ *function*), 64
 vec::sA (C++ *function*), 64
 vec::sB (C++ *function*), 64
 vec::sC (C++ *function*), 64
 vec::sD (C++ *function*), 64
 vec::sE (C++ *function*), 64
 vec::sF (C++ *function*), 64
 vec::store (C++ *function*), 64
 vec::swizzle (C++ *function*), 64
 vec::vec (C++ *function*), 63
 vec::w_ACCESS (C++ *function*), 64
 vec::x_ACCESS (C++ *function*), 64
 vec::y_ACCESS (C++ *function*), 64
 vec::z_ACCESS (C++ *function*), 64
 vector_class (C++ *type*), 22

W

weak_ptr_class (C++ *type*), 22
 work-group, 70
 work-item, 70